

Give io_uring a Chance!

Kyoung (lseo@uwaterloo.ca)

2024-07-23 Tue

Outline

Asynchronous IO in Linux

io_uring concepts

Examples

Conclusion

Readiness vs Completion

epoll

- ▶ Since Linux 2.5
- ▶ Is the socket **ready** for IO?

AIO

- ▶ subsumed by ...

io_uring

- ▶ Since Linux 5.1
- ▶ Was the kernel able to **complete** the IO on the socket?

Basic Concepts

Submission Queue, Completion Queue

Ring Provided Buffers

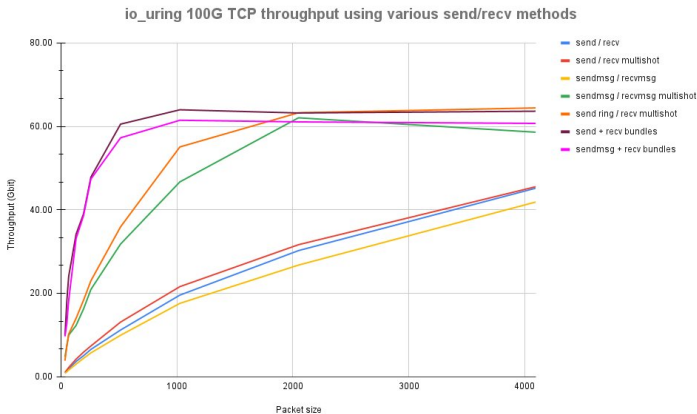
Multishot

Bundle

Sources of Potential Performance Gains

- ▶ “Batching”
 - - less poll rearming?
 - latency throughput tradeoff?
- ▶ Less system calls
- ▶ Less memcpy

proxy server

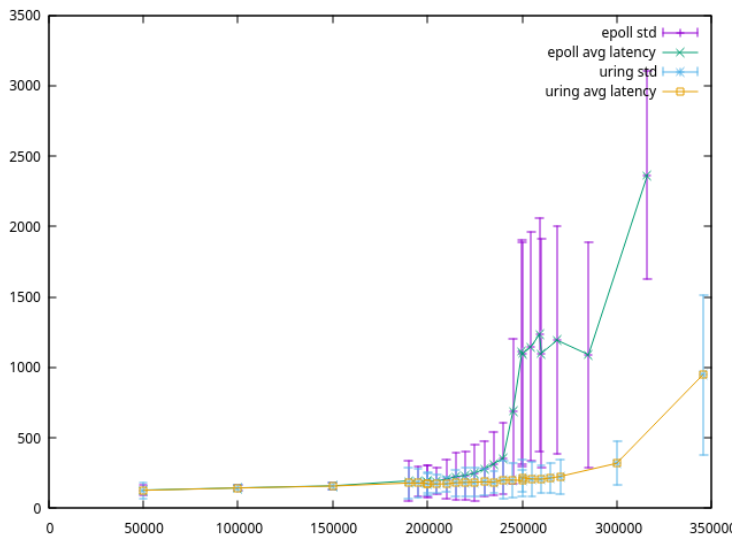


- ▶ multishot + provided buffers + bundles
- ▶ drastically reduced memory copies leading to high throughput


memcached

What would happen if we replace the epoll based event loop with io_uring ?

memcached



memcached



memcached

- ▶ 15% increase in throughput
 - ▶ 50% reduction in 99 percentile tail latency
- ... we are still coming up with explanations :D

Conclusion

Reasons to use `io_uring`

- ▶ Continuously being improved by very talented engineers at Meta
- ▶ Potential massive performance gain under certain scenarios

Challenges

- ▶ Lack of examples that demonstrate idioms
- ▶ Can be challenging to (re)design application logic and state machines

References

- ▶ <https://github.com/axboe/liburing/blob/master/examples/proxy.c>
- ▶ https://github.com/axboe/liburing/wiki/io_uring-and-networking-in-2023
- ▶ https://github.com/axboe/liburing/wiki/What's-new-with-io_uring-in-6.10
- ▶ https://git.uwaterloo.ca/lseo/memcached-io_uring