

Cpp
North
2025

C++ Contracts

a Meaningfully Viable Product

Andrei Zissu

C++ Contracts – a Meaningfully Viable Product

Andrei Zissu

Morphisec – preventive cyber security

- **WG21 member for the last 3 years**
- **Active in SG21**
 - **Not one of the main authors of the Contracts proposal (P2900)**

What This Talk Is About

- A moderately deep dive into contracts in C++26
- With some detail on considerations, controversies and paths not taken
- I'll strive to make the terminology easily accessible
 - Not always stick to the official terms
- I'll try not to focus on theory
 - But will sometimes present the important parts
 - Including relevant P2900 guiding principles, present in proper context
- I will be borrowing heavily from P2900 (the C++ contracts proposal, previously known as the contracts MVP) and P2899 (the rationale paper)
 - Occasionally other papers will also be mentioned and possibly borrowed from
- Not nearly enough time to cover everything I would have wanted to

Some quick terminology

Precondition Specifiers

```
1  auto div(auto x, auto y)
2      pre(y != 0)
3  {
4      return x/y;
5  }
6
7  int main()
8  {
9      return div(1, 0);
10 }
```

Output of x86-64 gcc (contracts natural syntax) (Compiler #1) ✎ ✕

A ▾ ☐ Wrap lines ☐ Select all

ASM generation compiler returned: 0

Execution build compiler returned: 0

Program returned: 139

```
contract violation in function div<int, int> at /app/example.cpp:2: y != 0
[assertion_kind: pre, semantic: enforce, mode: predicate_false, terminating: yes]
terminate called without an active exception
Program terminated with signal: SIGSEGV
```

- “Precondition” is a plain language term, while “precondition specifier” refers to a syntactic construct.

Postcondition Specifiers

```
1 auto plus(const unsigned x, unsigned y)
2     post(r: r >= x)
3 {
4     return x - y;
5 }
6
7 int main()
8 {
9     return plus(1, 1);
10 }
```

Output of x86-64 gcc (contracts natural syntax) (Compiler #1) ✎ ✕

A ▾ ☐ Wrap lines ☐ Select all

ASM generation compiler returned: 0

Execution build compiler returned: 0

Program returned: 139

contract violation in function plus at /app/example.cpp:2: r >= x

[assertion_kind: post, semantic: enforce, mode: predicate_false, terminating: yes]

terminate called without an active exception

Program terminated with signal: SIGSEGV

- The term “postcondition” is used in everyday language, whereas “postcondition specifier” denotes a specific syntactic element.
- *r* is directly naming the function return value – we’ll cover this later
- Also keep the const parameter in mind for later...

Function Contract Specifiers

- **The collective term for all pre and post specifiers (for preconditions and postconditions)**
- **Sometimes referred to in short just as “Contract Specifiers”**
- **I will often refer to them simply as pre and post**

Assertion Statements (a.k.a `contract_assert`)

```
1 auto plus(unsigned x, unsigned y)
2 {
3     auto ret = x - y;
4     contract_assert(ret >= x);
5     return ret;
6 }
7
8 int main()
9 {
10     return plus(1, 1);
11 }
```

Output of x86-64 gcc (contracts natural syntax) (Compiler #1) ✕

A ▾ ☐ Wrap lines ☒ Select all

ASM generation compiler returned: 0

Execution build compiler returned: 0

Program returned: 139

contract violation in function plus at /app/example.cpp:4: ret >= x

[assertion_kind: assert, semantic: enforce, mode: predicate_false, terminating: yes]

terminate called without an active exception

Program terminated with signal: SIGSEGV

Contract Assertions

- The collective term for the syntactic construct specified by function contract specifiers (pre and post) and assertion statements.
- Yes, this includes pre and post, while `contract_assert` is referred to only as “assertion statement”.
- For convenience I will be referring to them mostly as “CA” – short for “contract assertion” (which is a bit of a mouthful)

Contract Predicates

- The C++ code inside a contract assertion
- Must be contextually convertible to bool
- We can usually mix and match the terms “contract assert” and “contract predicate”
 - But occasionally they do have important distinguishing differences
 - Like *ignore* semantics

Some Bikeshedding History (or whatever you choose to name it)

Contract Assertions

- **Started off as CCA (Contract-Checking Annotation) - renamed since some committee members didn't think the term "annotations" is appropriate. Still being informally used (myself included).**
- **Also proposed: CAA (Contract-Assertion Annotation) - didn't catch on**
- **I proposed "contract clauses"**

contract_assert

- **Other languages use “assert”.**
- **But that would conflict with the existing assert macro.**
- **So pending a future solution, we needed another name.**
- **Over 40 alternatives were collected and explored in P2961R2 -
*A natural syntax for Contracts***

- **contractassert**
- **mustexpr**
- **dyn_assert**
- **musthold**
- **asrtexpr**
- **stdassert**
- **trueexpr**
- **co_assert**
- **ccassert** (this one is mine - disqualified as “too clever for most users”)
- **contract_assert**
- **std_assert**
- **dyn_check**
- **mustbetrue**
- **assertexpr**
- **assertion_check**
- **cppassert**
- **dynamic_assert**
- **cca_assert**
- **assrt**
- **runtime_assert**
- **_Assert**
- **xpct**
- **assert_check**
- **Assert2**

- **cpp_assert**
- **affirm**
- **__assert**
- **assess**
- **insist**
- **asrt**
- **cassert**
- **aver**
- **posit**
- **enforce**
- **audit**

- **claim**
- **ass**
- **must**
- **confirm**
- **assertion**
- **ensure**
- **chk**
- **verify**
- **expect**
- **check**

- **Finalists - `contract_assert` and `assertexpr`,**
- **`contract_assert` won the day.**
- **I voted against `assertexpr` as it resembles `constexpr`, which is used only in declarations. Given that `contract_assert` is currently a statement and not an expression (see later) I'm particularly glad we chose this name.**

Syntactic restrictions, and how they came to be

Multiple Declarations

– because we love IFNDR!

- (Function) contract specifiers (pre/post) must always be declared on function first-declarations
 - First declaration (used but not defined by the standard – until P2900 that is) - “a function from which no other declaration is reachable” - i.e. 1st declaration found by compiler in current TU (a.k.a .cpp file)
 - Otherwise it’s a “redeclaration” - that includes function body when defined separately from first declaration
 - Contracts repetition in function redeclarations is allowed (not mandatory, since compilers don’t really need them)
 - Mainly for user friendliness and to avoid further include hell

Multiple Declarations

– because we love IFNDR!

- **Ill-formed if repeated contract specifiers differ from reachable first declaration**
 - **Equivalence is determined by token equality (same as ODR) with renaming allowed for for function/template parameters and returned value**
 - **As long as the contract specifications do refer to the same entities**
- **Redeclarations in same TU cannot contain lambdas - would automatically render them “not the same” due to different type for each lambda**

Multiple Declarations

– because we love IFNDR!

- Well-formed renaming:

```
1  struct C
2  {
3      double div(double x)
4      |   pre(x != 0);
5  };
6
7  double C::div(double y)
8      pre(y != 0)
9  {
10     return y;
11 }
12
13 int main()
14 {
15     return C{}.div(1);
16 }
```

Multiple Declarations

– because we love IFNDR!

- Ill-formed renaming:

```
1 struct C
2 {
3     double div(double x, double y)
4     |
5     |   pre(y != 0);
6 }
7
8 double C::div(double x, double y)
9 |
10 |   pre(x != 0)
11 |
12 |   {
13 |       return x/y;
14 |   }
15
16 int main()
17 {
18     return C{}.div(1, 1);
19 }
```

Output of x86-64 gcc (contracts natural syntax) (Compiler #1) ✎ ✕

A ▾ ☐ Wrap lines ☐ Select all

<source>:8:11: error: mismatched contract condition in declaration

```
8 |   pre(x != 0)
  |           ^^^
```

<source>:4:15: note: previous contract here


```
4 |   pre(y != 0);
  |           ^^^
```

Compiler returned: 1

Semantics

Constification

– our favorite bowel non-movement

- officially called “implicit const-ness”
- often auto-corrected to “constipation” 
- Main motivation:
 - catch bugs like “assert(my_map[“universal_answer”] == 42)”
 - avoid contracts producing side effects, whether inadvertently or intentionally

Principle 6: No Destructive Side Effects

Contract assertions whose predicates, when evaluated, could affect the correctness of the program should not be supported.

- not entirely clear cut - e.g. can the mere extra clock cycles affect program correctness?

Constification

– our favorite bowel non-movement

- Effect:
 - Entities referenced by CAs (external to the CA itself) are treated similarly to data members in const member functions - with decltype still reporting the original constness
 - implicit const is shallow
 - objects pointed to by pointers are not constified
 - objects referred to by reference are constified
 - globals originally not included in constification - changed later to improve teachability

```
int global = 0;

void f(int x, int y, char *p, int& ref)
{
    pre((x = 0) == 0)           // error: assignment to const lvalue
    pre((*p = 5))               // OK
    pre((ref = 5))              // error: assignment to const lvalue
    pre((global = 2))           // error: assignment to const lvalue
    {
        int* gp = &global;
        contract_assert((gp = nullptr)); // error: assignment to const lvalue
        contract_assert((*gp = 6));      // OK
    }
}
```


Constification

– our favorite bowel non-movement

- controversial opt out - `const_cast`
 - must be used with utmost care, as the outcome is UB if the original object is `const`

```
int g(int i, const int j)
    pre(++const_cast<int&>(i))    // OK (but discouraged)
    pre(++const_cast<int&>(j))    // undefined behavior
    post(++const_cast<int&>(i))   // OK (but discouraged)
    post(++const_cast<int&>(j))   // undefined behavior
{
    int k = 0;
    const int l = 1;
    contract_assert(++const_cast<int&>(k)); // OK (but discouraged)
    contract_assert(++const_cast<int&>(l)); // undefined behavior
}
```

Constification

– our favorite bowel non-movement

- **Objections:**
 - **modifies overload resolution, therefore contracts may actually be evaluating different code**

```
struct X {};  
bool p(X&) { return true; }  
bool p(const X&) { return false; }
```

```
void my_assert(bool b) { if (!b) std::terminate(); }
```

```
void f(X x1)  
    pre(p(x1))           // fails  
{  
    my_assert(p(x1));    // passes  
  
    X x2;  
    contract_assert(p(x2)); // fails  
    my_assert(p(x2));      // passes  
}
```

Constification

– our favorite bowel non-movement

- **Objections:**
 - **modifies overload resolution, therefore contracts may actually be evaluating different code**
 - **counter - such overloads are bugs on their own**
 - **counter counter - contracts must be useful with real world code, not educate programmers**
 - **counter counter counter – on the contrary, contracts should encourage correct programming practices**
 - **counter - side effects in contracts create a different program – rendering such contracts potentially meaningless/harmful**
 - **counter - may help discover latent side effects in existing assert statements**

Constification

– our favorite bowel non-movement

- **Objections:**
 - **shallow const**
 - **constification is incomplete and inconsistent**
 - **in particular, pointers and references are inconsistent with each other**
 - **counter - pointer/reference inconsistency is an existing language feature**
 - **counter - automatic deep const would be very complex to specify and implement**

Constification

– our favorite bowel non-movement

- Objections:
 - constification blocks some common non-const usages, such as logging and `std::map::operator[]`
 - counter - `std::map::operator[]` is bad, `std::map::at()` should be used instead
 - no such easy solution with things like logging though
 - counter - `const_cast` as a controversial escape hatch
 - other proposed escape hatches:
 - `mutable` (applied to CA),
 - `operator noconst` (applied to expression in CA):

```
// Legacy API, check_valid doesn't modify x but isn't const-qualified
struct X;
bool check_valid(X& x);

void f(X& x)
    pre(check_valid(x)); // Ill-formed due to implicit const on x

void g(X& x)
    pre(noconst(x).is_valid()); // OK: noconst(x) treats x as non-const for this expr
```

Postconditions

Postconditions- Referring to the Result Object

– what is a name?



- Post may specify a name for the return value, valid only within the post itself
- That name will capture the result object “on the fly”
 - Even when the function returns an unnamed temporary - no other way in the language to do this for all objects
 - Although it is feasible for user-provided constructors, but only where those are available

Postconditions - Referring to the Result Object – what is a name?

- Similar to references but not really one, also compared to structured bindings
 - Main difference is that `decltype` doesn't see it as a reference

```
1  #include <type_traits>
2
3  int f() post(r: std::is_reference_v<decltype(r)>)
4  {
5      int x;
6      int& ref = x;
7      static_assert(std::is_reference_v<decltype(ref)>);
8      return 42;
9  }
10
11 int main()
12 {
13     f();
14     return 0;
15 }
```

Output of x86-64 gcc (contracts natural syntax) (Compiler #1) [↗](#) [✕](#)

 ☐ Wrap lines  Select all

ASM generation compiler returned: 0

Execution build compiler returned: 0

Program returned: 139

```
contract violation in function f at /app/example.cpp:3: std::is_reference_v<decltype(r)>
[assertion_kind: post, semantic: enforce, mode: predicate_false, terminating: yes]
terminate called without an active exception
Program terminated with signal: SIGSEGV
```


Postconditions- Referring to the Result Object

– what is a name?

```
struct S {  
    S();  
    S(const S&) = delete; // noncopyable, nonmovable  
    int i = 0;  
    bool foo() const;  
};  
  
const S f()  
{  
    post(r: const_cast<S&>(r).i = 1) // OK (but discouraged)  
    {  
        return S{};  
    }  
}  
  
const S y = f(); // well-defined behavior  
bool b = f().foo(); // well-defined behavior
```

- **const_cast**, although discouraged, is actually safe in this case - result object being captured on the fly is not yet assigned to the call site target variable, therefore it cannot yet be const

Referring to Parameters in Postconditions – wait, isn't const meaningless when passing by value?

- TLDR: by-value function parameters odr-used (a.k.a consumed) in a post CA - must be const, in all the function declarations

```
void f(int i) post ( i != 0 );           // error: i must be const.
```

```
void g(const int i) post ( i != 0 );  
void g(int i) {}                        // error: missing const for i in definition
```

```
void h(const int i) post (i != 0);  
void h(const int i) {}  
void h(int i);                          // error: missing const for i in redeclaration
```

Referring to Parameters in Postconditions – wait, isn't const meaningless when passing by value?

- **Motivation (also in P2521 - 3.13):**
 - **CAs refer to parameters directly (not to copies thereof)**
 - **post CAs are intended to work on initial parameter values**
 - **contracts impossible to reason about if parameters may mutate**
- **const must be consistently applied in all the function declarations**
- **with future contract capture clauses (not in C++26) function parameters won't have to be const – since intent will be explicit**

Evaluation and Contract-Violation Handling

Point of Evaluation

– who said it's a single point?

function call

init params |

pre

eval func body

contract_assert (when reached by control flow)

init result object (this means actual caller context object for non-trivially copyable returned types)

destruct local vars

post

destruct params → bind result object

- pre/post see function parameters (or possibly their register copies) but not local variables
- result object is initialized before post (thus post can see it), but not yet considered bound
 - caller context const (if present) doesn't apply yet (meaning const_cast of result object in post is safe from UB)
- post is only evaluated on normal exit
 - evaluation when exiting due to exception may be possible in the future
- post is evaluated only after local variables destruction - since those might influence what post is checking

Evaluation Semantics and selection thereof

– because nothing is pre-ordained

- **checking semantics**
 - **observe**
 - **terminating semantics**
 - **enforce**
 - **quick_enforce** (was “Louis semantic”, prior to yet another bikeshedding procedure)
- **ignore semantic**
 - **unlike C assert, it odr-uses and must be well-formed (which prevents code rot)**

Evaluation Semantics and selection thereof

– because nothing is pre-ordained

- upon violation
 - at run time
 - observe
 - invoke violation handler
 - continue execution (upon normal return)
 - enforce
 - invoke violation handler
 - terminate (upon normal return)
 - quick_enforce
 - terminate immediately
 - at compile time
 - observe – issue diagnostic (a.k.a warning)
 - enforce and quick_enforce – make program ill-formed
 - ignore – doing nothing in both cases (other than odr-use)

Evaluation Semantics and selection thereof

– because nothing is pre-ordained

- **selection of contract semantics**
 - **implementation-defined - because different implementations/platforms have different needs**
 - **no specification of when this happens**
 - **any build stage**
 - **run time**
 - **allowing dynamic contracts configuration**
 - **build modes (e.g. debug and release) are no longer required for this**
 - **but of course are still allowed**
 - **different semantics may be selected for different CAs in the same TU**
 - **or even for the same CA at different evaluations!**

Evaluation Semantics and selection thereof – because nothing is pre-ordained

- compliant implementations may offer any non-empty subset of the 4 semantics
 - meaning that ignore semantic alone is compliant, which entails:
 - enforcing CAs being well-formed
 - odr-use
 - side effects possible, e.g. initializing static data members of class templates
- chosen semantic explicitly undetectable from code at compile time
 - to avoid contracts changing observed behavior

Principle 3: Concepts Do Not See Contracts

The mere presence of a contract assertion on a function or in a block of code should not change the satisfiability of a concept, the result of overload resolution and SFINAE, the branch selected by `if constexpr`, or the value returned by the `noexcept` operator.

Elision and Duplication - I swear to perhaps check the truth, and check it again and again and again, so help me the semantic

- **any CA may be evaluated 0...N times**
 - **not necessarily with the same semantics**
- **repetition is required for supporting different caller and callee contract semantics**
 - **likely to affect mainly pre and post**
 - **up to 2 evaluations should normally suffice,**
 - **other than multiple repetitions to test for unwanted CA side effects**
- **Implementation requirements**
 - **define an upper bound for repetition**
- **Implementation recommendations**
 - **allow users to configure any number of repetitions**
 - **make the default one single evaluation without repetitions**

Elision and Duplication - I swear to perhaps check the truth, and check it again and again and again, so help me the semantic

- Elision – as in skip
 - Only if compiler can prove that predicate:
 - Returns true – typically if any of the following holds:
 - It's evaluated at compile time
 - It's guaranteed by previous CAs (in checked terminating semantics)
 - Example:
`contract_assert(x>1);`
`contract_assert(x>0);` `// may be elided`
 - Cannot throw, longjmp or terminate

Elision and Duplication - I swear to perhaps check the truth, and check it again and again and again, so help me the semantic

- Elision – as in rephrase
 - The compiler may also generate an equivalent expression
 - Only required to cover the well-defined cases
 - UB in original predicate is fair game
 - Existing side effects may be ditched in the process – all or nothing
 - No new side effects may be introduced

Principle 6: No Destructive Side Effects

Contract assertions whose predicates, when evaluated, could affect the correctness of the program should not be supported.

```
int i = 0;
void f() pre ((++i, true));
void g() {
    f(); // i may be 0, 1, 17, etc.
}
```

The Contract-Violation Handler

– you can't avoid termination (or can you?)

- **definition**
 - a function named `::handle_contract_violation`
 - single argument of type `const std::contracts::contract_violation&`
 - returns `void`
 - may be `noexcept` (but definitely doesn't have to be)

The Contract-Violation Handler

– you can't avoid termination (or can you?)

```
class contract_violation
{ // no user-accessible constructor; cannot be copied, moved, or assigned to
public:
    const char*          comment() const noexcept;
    contracts::detection_mode detection_mode() const noexcept; // predicate_false, evaluation_exception
    exception_ptr         evaluation_exception() const noexcept;
    bool is_terminating() const noexcept;
    assertion_kind        kind() const noexcept;           // pre, post, assert
    source_location        location() const noexcept;
    evaluation_semantic semantic() const noexcept; // (ignore), observe, enforce, (quick_enforce)
};
```

The Contract-Violation Handler

– you can't avoid termination (or can you?)

- properties
 - implementations must provide default handlers
 - recommended to output the *contract_violation* info and to be noexcept
 - default handler cannot be directly called by user code
 - but it can indirectly
 - *void invoke_default_contract_violation_handler(const contract_violation&);*};
 - but only from custom violation handlers, since
 - users have no way of creating *contract_violation* objects

The Contract-Violation Handler

– you can't avoid termination (or can you?)

- **properties**
 - **implementations may allow the handler to be replaceable**
 - **i.e. provide a function with the same name, parameter types and return type**
 - **which for replaceable functions does not result in linker errors (ambiguous symbol)**
 - **may have a different exception specification than the default handler**
 - **hence the standard library provides only an implementation but not an includable declaration of the default handler**
 - **same as global operators new and delete**
 - **some platforms/implementers may consider that an unacceptable security risk – that's ok**

The Contract-Violation Handler

– you can't avoid termination (or can you?)

- **libraries control what contracts they contain**
- **application controls how contracts are handled**
 - **with the custom handler instrumental in their toolbox**
- **installed at link time , because:**
 - **different TUs and libs may be compiled at different times and with different toolchains**
 - **too risky security-wise at run time**

The Contract-Violation Handler

– you can't avoid termination (or can you?)

- replacing the handler doesn't require recompiling the whole application (and libraries)
 - otherwise it would be a major impediment to contracts adoption
- *contract_violation* object requirements
 - may only be produced by the implementation (not by user code)
 - may reside anywhere other than the heap
 - can be accessed by violation handler (inc. other objects pointed to) without lifetime concerns
- any exception in the CA itself is caught and treated by the handler as an additional detection mode

Throwing Violation Handlers and `contract_assert(false)` – are you calling me a liar?


- Contract violation handlers are allowed to return via throw or longjmp
- Circumventing enforce semantic - a feature, not a bug!
- Bjarne himself required this in P2698R0
 - Stating that some systems cannot afford termination
- “Terminating semantics” may be a bit of a misnomer
 - Their actual guarantee (*enforce* especially) is of not allowing execution of the immediately following code
 - Most likely buggy code following a failed CA will probably crash anyway
 - Or trigger any other manifestation of UB

Throwing Violation Handlers and `contract_assert(false)` – are you calling me a liar?

- What does `noexcept(contract_assert(false))` return?

Principle 1: Prime Directive

The presence or evaluation of a contract assertion in a program should not alter the correctness of that program (i.e., the property that evaluation of the program does not violate any provisions of its plain-language contract).



- Having contracts alter an exception specification would violate the prime directive
- But a previously non-throwing function/expression may now be throwing
 - Depending on whether the installed violation handler is `noexcept`, which is unknown at compile time
 - For which reason it must be assumed to be true (on platforms supporting throwing violation handlers)

Throwing Violation Handlers and `contract_assert(false)` – are you calling me a liar?

- Rock and hard place
 - Tell the truth and violate the prime directive
 - Or lie and be damned for all eternity
 - And we did fight over this almost for all eternity...
 - And as in other similar cases we chose to avoid the land mine altogether

Principle 14: Choose Ill-Formed to Enable Flexible Evolution

When no clear consensus has become apparent regarding the proper solution to a problem that Contracts could address, the relevant constructs are left ill-formed.

Throwing Violation Handlers and `contract_assert(false)` – are you calling me a liar?

- Rock and hard place
 - But rather with a little cheating
 - We made `contract_assert` a statement rather than an expression
 - Well, I lied a bit too - this does make `noexcept(contract_assert(false))` ill-formed...
 - Now `contract_assert` is slightly more limited than C `assert` - can't be a sub-expression
 - Which immediately invoked lambdas help mitigate
 - Instead of this:
 - `const int j = (contract_assert(i > 0), i);`
 - We can write this:
 - `const int j = ([i]{ contract_assert(i > 0); }(), i);`
 - Which works since this works:
 - `const int j = (void{}, i);`

Final Word

- **Contracts are complicated**
 - **As is C++**
- **Which is why it took so long to have them delivered**
- **But you won't normally care about all these details**
 - **Like about most other C++ details**
- **But you can leverage them when necessary**
 - **You already know what I'm going to say here...**
- **At the end of the day, contracts are a great tool**
 - **Like C++, you know the drill...**

Final Word

- **Many thanks to the many contributors who made Contracts finally happen in C++26!**
- **And special thanks to 2 of them who also thoroughly reviewed my presentation draft: Timur Doumler and Joshua Berne**
 - **Which makes any errors in this presentation definitely only my own**

Thank You!

Let's get in touch!

andrziss@gmail.com

<https://www.linkedin.com/in/andreizissu/>

Questions?