

Cpp
North
2025

Heaps Don't Lie

Guidelines for Memory Allocation in C++

Mathieu Ropert

Heaps Don't Lie

Guidelines for Memory Allocation in C++



Don't use malloc()



“

Don't use *malloc()* while the
plane is in the air

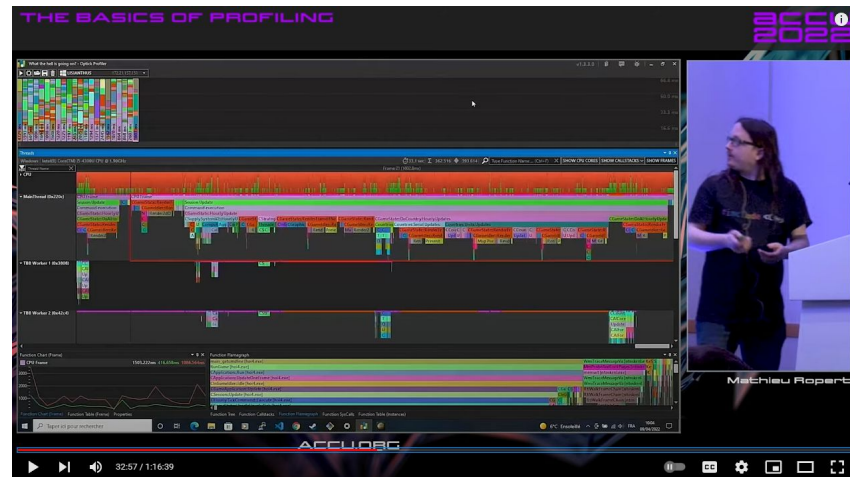


“



On performance

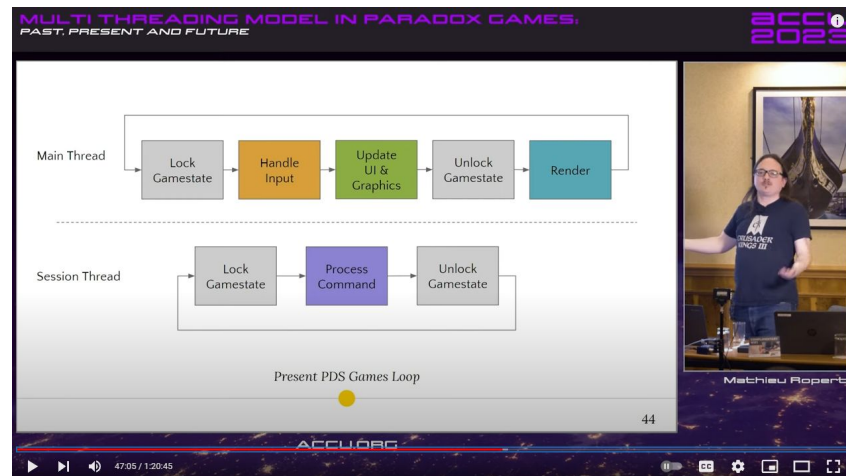
- Profiling





On performance

- Profiling
- Multithreading





On performance

- Profiling
- Multithreading
- Caches & data locality

DATA ORIENTED DESIGN AND ENTITY COMPONENT SYSTEM EXPLAINED

Video Sponsored By think-cell

ACCU conference 2024

Bloomberg

Mathieu Ropert

L1 Cache

L2 Cache

CPU vs RAM

21

ACCU.ORG

16:23 / 1:21:22



On performance

- Profiling
- Multithreading
- Caches & data locality
- Allocations

this

Hello!



I am **Mathieu Ropert**

I'm a C++ programmer, game developer and meetup organizer. I'm available for consulting and training!

You can reach me at:



mro@puchiko.net



[@MatRopert](https://twitter.com/MatRopert)



[@matropert.bsky.social](https://bsky.app/profile/matropert.bsky.social)



<https://mropert.github.io>



About this talk

- How does the heap work?
- 10 guidelines for handling memory allocations
- For beginners and experts alike

1

What's in a `malloc()`?

Quite a bit, it turns out



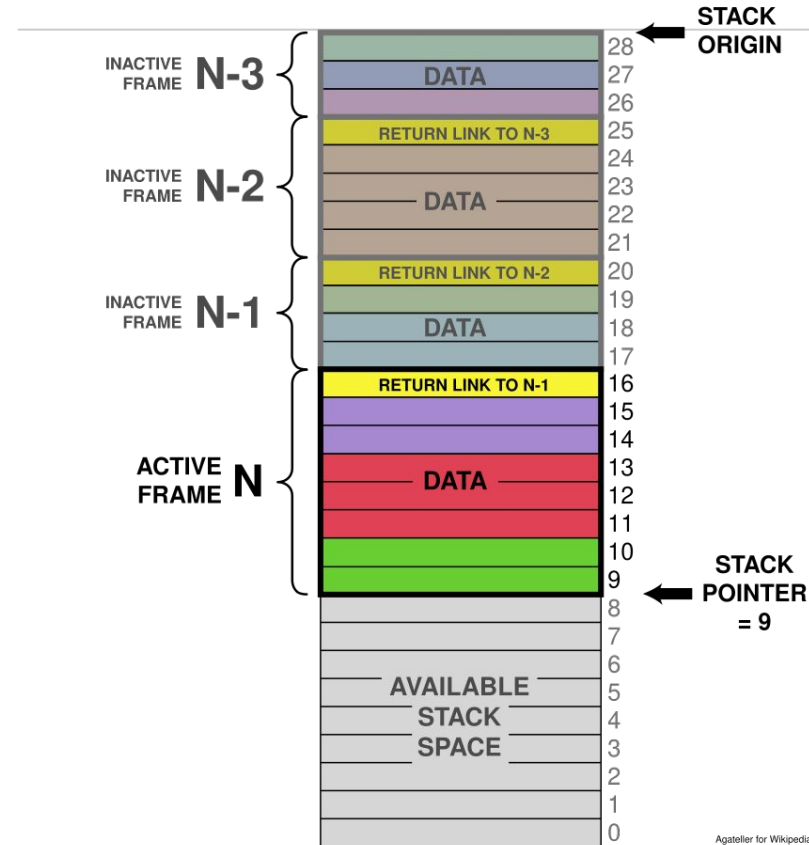
Defining “allocation”

- The act of reserving *uninitialized* bytes
- Not concerned with zeroing or constructing objects in it
- Likewise for deallocation



A (stack) frame of reference

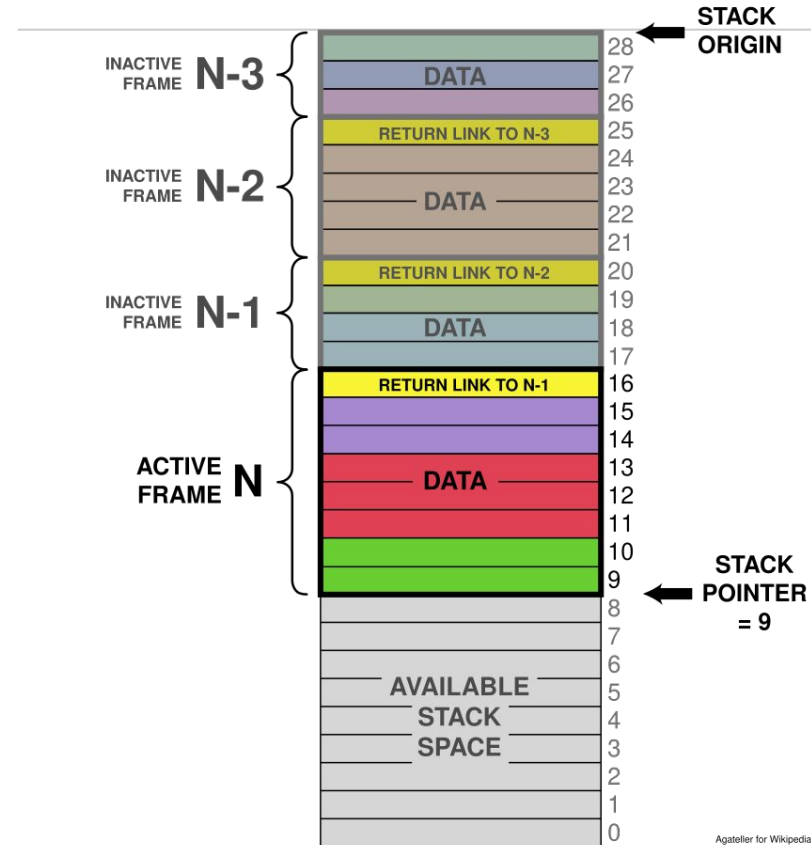
- Stack allocation
- Increment stack pointer and you're done
- One register add
- No library or operating system call





Stack memory limitations

- Size of allocation defined at compile time
- Limited in space
- Automatically reclaimed on scope exit



*Guideline #1: prefer stack alloc
when possible*



“



Enter malloc()


- Can allocate arbitrary amounts of memory
- Lives until matching `free()` is called
- Can mimic stack lifetime behaviour with `unique_ptr` and RAI




malloc() process



- Try to get a free block of size N from free list



- Grab a free page, initialize it as page of size objects



- Use an OS system call to get more pages
- Swap some physical pages to disk

*Guideline #2: malloc() is usually
fast, except when it isn't*



“



`malloc()` **implementation**

- Free blocks and pages are managed by C library implementation
- Virtual memory and paging depend on OS
- First one can be replaced by user-provided implementation



The slide about **MSVC**

- The default allocator that comes with MSVCRT is not great
- Locks mutex on `free(NULL)` in Debug builds
- Bad performance when allocating/freeing in parallel



The slide about MSVC



ree

*Guideline #3: prefer third party
heap implementation on Windows*



“

2

Allocation **best** practices

Experiences collected over the years



The recurring pathfinder test

- Somewhat simple exercise in principle
- Seen dozens of implementations over the years
- Some allocation related mistakes keep repeating

THE BASICS OF PROFILING

A* Refresher

- Open set: nodes/square reachable but not explored
- Closed set: nodes/squares fully explored
- Pick best candidate in open set, add neighbours to open set, repeat until destination is reached

40

Mathieu Robert



Exploring neighbours

- For a given cell (x, y)
- Check each neighbours
- Filter out/return the ones that are within map bounds and are traversable



Exploring neighbours

```
std::vector<Point> find_neighbours( Point p ) const
{
    std::vector<Point> neighbours;
    if ( p.x_ > 0 ) neighbours.emplace_back( p.x_ - 1, p.y_ );
    if ( p.x_ < size_.x_ - 1 ) neighbours.emplace_back( p.x_ + 1, p.y_ );
    if ( p.y_ > 0 ) neighbours.emplace_back( p.x_, p.y_ - 1 );
    if ( p.y_ < size_.y_ - 1 ) neighbours.emplace_back( p.x_, p.y_ + 1 );
    return neighbours;
}
```

Allocates!

Allocates!

Allocates!



Exploring neighbours

```
std::vector<Point> find_neighbours( Point p ) const
{
    std::vector<Point> neighbours;
    neighbours.reserve( 4 );
    if ( p.x_ > 0 ) neighbours.emplace_back( p.x_ - 1, p.y_ );
    if ( p.x_ < size_.x_ - 1 ) neighbours.emplace_back( p.x_ + 1, p.y_ );
    if ( p.y_ > 0 ) neighbours.emplace_back( p.x_, p.y_ - 1 );
    if ( p.y_ < size_.y_ - 1 ) neighbours.emplace_back( p.x_, p.y_ - 1 );
    return neighbours;
}
```

Allocates!

*Guideline #4: reserve final
container size when known rather
than rely on geometric growth*


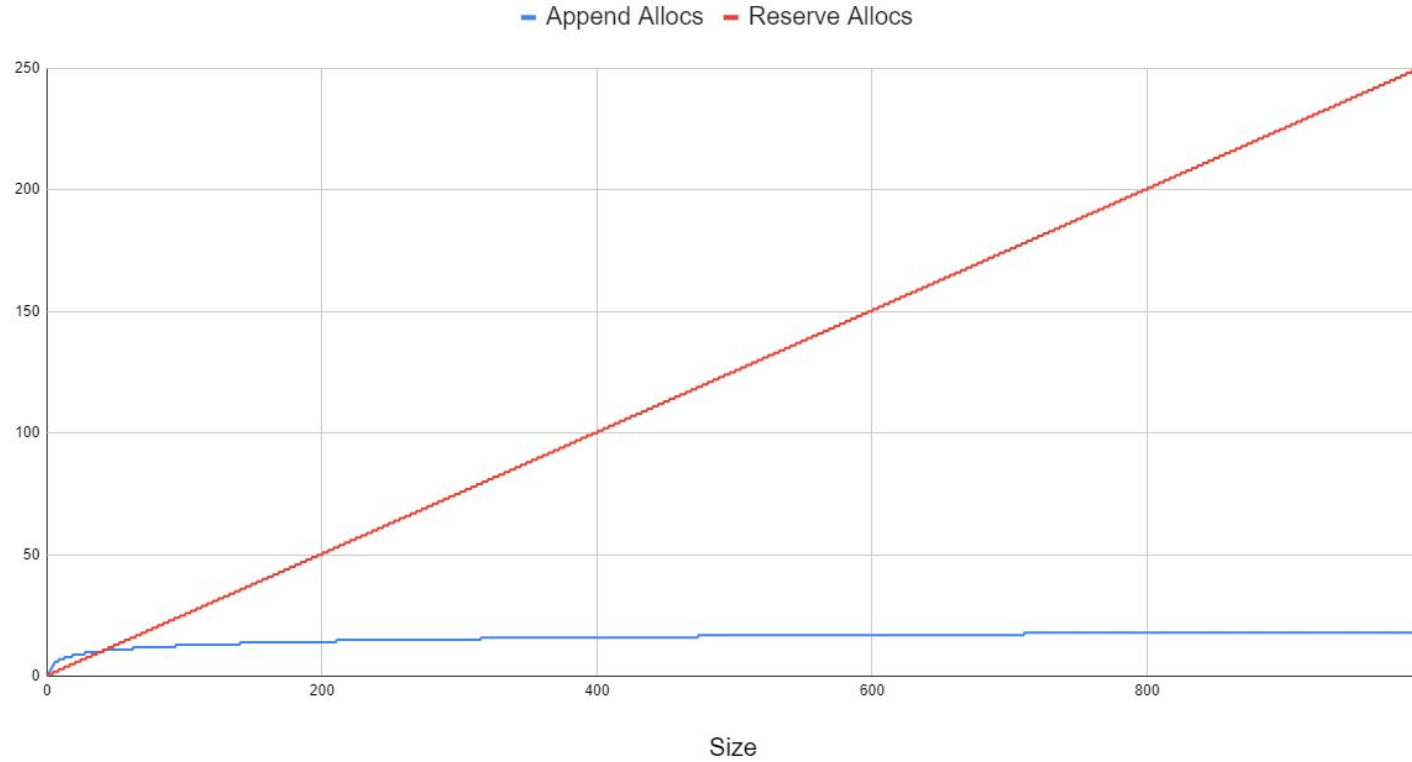
“



Abusing reserve()

```
void add_neighbours( const std::vector<Point>& added )  
{  
Allocates! neighbours_.reserve( neighbours_.size() + added.size() );  
std::copy( begin( added ), end( added ),  
           std::back_inserter( neighbours_ ) );  
}
```

Append Allocs vs Reserve Allocs



*Guideline #5: avoid using
reserve() with constant delta in
loops*



“



Zooming out

```
while ( !search_queue.empty() )
{
    const Point current = search_queue.top();
    search_queue.pop();
    Allocates! for ( const Point p : find_neighbours( current ) )
    {
        ...
    }
}
```




Exploring neighbours, **redux**

```
void find_neighbours( Point p, std::vector<Point>& neighbours ) const
{
    if ( p.x_ > 0 ) neighbours.emplace_back( p.x_ - 1, p.y_ );
    if ( p.x_ < size_.x_ - 1 ) neighbours.emplace_back( p.x_ + 1, p.y_ );
    if ( p.y_ > 0 ) neighbours.emplace_back( p.x_, p.y_ - 1 );
    if ( p.y_ < size_.y_ - 1 ) neighbours.emplace_back( p.x_, p.y_ + 1 );
}
```



Outer loop, revised

```
std::vector<Point> neighbours;
while ( !search_queue.empty() )
{
    const Point current = search_queue.top();
    search_queue.pop();
    neighbours.clear();
    find_neighbours( current, neighbours );
    for ( const Point p : neighbours ) { ... }
}
```

*Allocates!
(once)*

*Guideline #6: prefer output
parameters to returns for
containers*



“

*Guideline #7: reuse previous
allocations when possible*



“



Clearing structures

```
struct Frame {  
    std::vector<const Entity*> visible_entities_  
    std::vector<const Entity*> audible_entities_  
    std::string debug_message_  
    ...  
    void clear()  
    {  
        *this = Frame(); Frees!  
    }  
};
```



Reassigning containers

```
vector& operator=( vector&& other )  
{  
    clear();  
    shrink_to_fit();  
    swap( other );  
    return *this;  
};
```

Frees!



Reassigning containers

```
vector& operator=( vector&& other )
{
    clear();
    if ( !other.empty() )
    {
        shrink_to_fit();
        swap( other );
    }
    return *this;
};
```



Reassigning containers

```
vector& operator=( vector&& other )
{
    clear();
    if ( !other.empty()
        || ( other.empty() && capacity() < other.capacity() ) )
    {
        shrink_to_fit();
        swap( other );
    }
    return *this;
};
```




Clearing structures compared

```
void clear()
{
    *this = Frame();
}
```

```
void clear()
{
    visible_entities_.clear();
    audible_entities_.clear();
    debug_message_.clear();
}
```



Clearing structures compared

```
void render()  
{  
  Oops! frame_data_ = Frame();  
  ...  
}
```

```
void render()  
{  
  frame_data_.clear();  
  ...  
}
```

Guideline #8: prefer memberwise
.*clear()* over assignment to
empty struct

“

3

Custom allocators

And other alternatives



General purpose **allocator**

- Has to be “good enough” for everyone
- Can't make assumptions about usage patterns
- Compromises have to be made
- What else is there?



What about `alloca()`?

- Not standard (but very common)
- Dynamic stack allocation
- C API, need wrapper to add RAI back
- Fast



Drawbacks of `alloca()`

- Bound by stack size
- Limited space
 - Especially on threads
 - Especially on macOS
- Not all threading libraries support setting stack size (standard C++ doesn't)

Guideline #9: avoid using
alloca() because it's likely to
overflow the stack




“



Custom allocators

- Special-purpose allocators
- Found in various libraries and in the standard since C++17
- Drop some malloc() constraints for speed (and/or fragmentation)

Pool Allocator - Reclamation





Monotonic allocators

- Allocates by increasing a pointer until it runs out
- As fast as stack alloc
- Cannot deallocate one specific object
- Free all or nothing



Simple monotonic allocator

```
struct monotonic_buffer
{
    std::unique_ptr<std::byte[]> buffer_;
    std::byte *current_;

    explicit monotonic_buffer( std::size_t size )
        : buffer_( std::make_unique<std::byte[]>( size ) )
        , current_( buffer_.get() ) { }

    ...
};
```



Simple monotonic allocator

```
void* allocate( std::size_t size )
{
    constexpr auto align_mask = sizeof( std::max_align_t ) - 1;
    const auto alloc_size = ( size + align_mask ) & ~align_mask ;
    void* allocated = current_;
    current_ += alloc_size;
    return allocated;
}
```



Simple monotonic allocator

```
void deallocate( void* ptr ) { }
```

```
void deallocate_all()  
{  
    current_ = buffer_.get();  
}
```



Monotonic allocators

- Good for scoped allocations that can be budgeted
- Game levels
- State machines
- Operation, transaction and request processing

*Guideline #10: consider
monotonic allocators when
resources can be scoped and
budgeted*



“

4

Wrapping up

Would you like to know more?



In conclusion

- Dynamic allocation comes at an unpredictable runtime cost
- Size-up allocations once and reuse containers as much as possible
- Consider monotonic allocators, but beware of going further

Furthermore

“

*Furthermore, I think your build
should be destroyed*

“




Thanks!

Any **questions** ?

You can reach me at

 mro@puchiko.net

 @MatRopert  @matropert.bsky.social

 @mropert

 <https://mropert.github.io>