



MULTI-DIMENSIONAL

Bryce Adelstein Lelbach

Principal Architect



@blelbach

Standard C++ Library Evolution Chair, US Programming Languages Chair

Today, C++ has no reasonable standard abstraction for multi-dimensional data.

Today, C++ has no reasonable standard abstraction for multi-dimensional data.

C++ needs abstractions that allow us to write generic multi-dimensional code.

We want to write generic multi-dimensional code that is:

- Storage agnostic.
- Rank agnostic.
- Layout agnostic.
- Iteration pattern agnostic.
- Composable.

Today, C++ has no reasonable standard abstraction for multi-dimensional data.

The solution is coming in C++23:

`std::mdspan`

`std::mdspan`

- Non-owning; pointer + metadata.

`std::mdspan`

- Non-owning; pointer + metadata.
- Metadata can be dynamic or static.

`std::mdspan`

- Non-owning; pointer + metadata.
- Metadata can be dynamic or static.
- Parameterizes layout.

`std::mdspan`

- Non-owning; pointer + metadata.
- Metadata can be dynamic or static.
- Parameterizes layout and access.

```
template <typename Index, Index... Extents>  
class std::extents;
```

```
template <typename Index, Index... Extents>  
class std::extents;
```

```
std::extents e0{16, 32};
```

```
// Equivalent to:
```

```
std::extents<std::dynamic_extent, std::dynamic_extent> e1{16, 32};
```

```
e0.rank() == 2
```

```
e0.extent(0) == 16
```

```
e0.extent(1) == 32
```

```
template <typename Index, Index... Extents>  
class std::extents;
```

```
std::extents e0{16, 32};
```

```
// Equivalent to:
```

```
std::extents<std::dynamic_extent, std::dynamic_extent> e1{16, 32};  
std::dextents<2> e2{16, 32};
```

```
e0.rank() == 2
```

```
e0.extent(0) == 16
```

```
e0.extent(1) == 32
```

```
template <typename Index, Index... Extents>  
class std::extents;
```

```
std::extents e0{16, 32};
```

```
// Equivalent to:
```

```
std::extents<std::dynamic_extent, std::dynamic_extent> e1{16, 32};  
std::dextents<2> e2{16, 32};
```

```
e0.rank()      == 2  
e0.extent(0)   == 16  
e0.extent(1)   == 32
```

```
std::extents<16, 32> e3;
```

```
template <typename Index, Index... Extents>  
class std::extents;
```

```
std::extents e0{16, 32};
```

```
// Equivalent to:
```

```
std::extents<std::dynamic_extent, std::dynamic_extent> e1{16, 32};  
std::dextents<2> e2{16, 32};
```

```
e0.rank()      == 2  
e0.extent(0)   == 16  
e0.extent(1)   == 32
```

```
std::extents<16, 32> e3;
```

```
std::extents<16, std::dynamic_extent> e4{32};
```

```
template <typename Index, Index... Extents>  
class std::extents;
```

```
std::extents e0{16, 32};
```

```
// Equivalent to:
```

```
std::extents<std::dynamic_extent, std::dynamic_extent> e1{16, 32};  
std::dextents<2> e2{16, 32};
```

```
e0.rank()      == 2  
e0.extent(0)   == 16  
e0.extent(1)   == 32
```

```
std::extents<16, 32> e3;
```

```
std::extents<16, std::dynamic_extent> e4{32};
```

```
std::extents e5{16, 32, 48, 4};
```

```
template <
```

```
>
```

```
class std::mdspan;
```



```
template <class I,
```

```
>
```

```
class std::mdspan;
```

```
template <class I,  
          class Extents,
```

>

```
class std::mdspan;
```

```
template <class I,  
         class Extents,  
         class LayoutPolicy = std::layout right,
```

>

```
class std::mdspan;
```

```
template <class I,  
         class Extents,  
         class LayoutPolicy = std::layout right,  
         class AccessorPolicy = std::default accessor<T>>  
class std::mdspan;
```

```
template <class I,  
         class Extents,  
         class LayoutPolicy = std::layout right,  
         class AccessorPolicy = std::default accessor<T>>  
class std::mdspan;  
  
std::mdspan m0{data, 16, 32};  
// Equivalent to:  
std::mdspan<double, std::dextents<2>> m1{data, 16, 32};
```

```

template <class I,
          class Extents,
          class LayoutPolicy = std::layout right,
          class AccessorPolicy = std::default accessor<T>>
class std::mdspan;

std::mdspan m0{data, 16, 32};
// Equivalent to:
std::mdspan<double, std::dextents<2>> m1{data, 16, 32};

m0[i, j] == data[i * M + j]

```

```

template <class I,
          class Extents,
          class LayoutPolicy = std::layout right,
          class AccessorPolicy = std::default_accessor<T>>
class std::mdspan;

std::mdspan m0{data, 16, 32};
// Equivalent to:
std::mdspan<double, std::dextents<2>> m1{data, 16, 32};

m0[i, j] == data[i * M + j]

std::mdspan m2{data, std::extents<16, 32>{}};
// Equivalent to:
std::mdspan<double, std::extents<16, 32>> m3{data};

std::mdspan m4{data, std::extents<16, std::dynamic_extent>{32}};

```

Row-Major AKA Right

- C++, NumPy (default)
- Rightmost extent is contiguous

```
mdspan A{data, N, M};  
mdspan A{data, layout_right::mapping{N, M}};
```

```
A[i, j]      == data[i * M + j]  
A.stride(0) == M  
A.stride(1) == 1
```


Row-Major AKA Right

- C++, NumPy (default)
- Rightmost extent is contiguous

```
mdspan A{data, N, M};  
mdspan A{data, layout_right::mapping{N, M}};
```

```
A[i, j] == data[i * M + j]  
A.stride(0) == M  
A.stride(1) == 1
```

Location	Element
0	a_{11}
1	a_{12}
2	a_{21}
3	a_{22}

$$\begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}$$

Row-Major AKA Right

- C++, NumPy (default)
- Rightmost extent is contiguous

Column-Major AKA Left

- Fortran, MATLAB
- Leftmost extent is contiguous

```
mdspan A{data, N, M};  
mdspan A{data, layout right::mapping{N, M}};
```

```
A[i, j] == data[i * M + j]  
A.stride(0) == M  
A.stride(1) == 1
```

```
mdspan B{data, layout left::mapping{N, M}};
```

```
B[i, j] == data[i + j * N]  
B.stride(0) == 1  
B.stride(1) == N
```

Location	Element
0	a_{11}
1	a_{12}
2	a_{21}
3	a_{22}

$$\begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}$$

Row-Major AKA Right

- C++, NumPy (default)
- Rightmost extent is contiguous

```
mdspan A{data, N, M};  
mdspan A{data, layout right::mapping{N, M}};
```

```
A[i, j] == data[i * M + j]  
A.stride(0) == M  
A.stride(1) == 1
```

Location	Element
0	a_{11}
1	a_{12}
2	a_{21}
3	a_{22}

$$\begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}$$

Column-Major AKA Left

- Fortran, MATLAB
- Leftmost extent is contiguous

```
mdspan B{data, layout left::mapping{N, M}};
```

```
B[i, j] == data[i + j * N]  
B.stride(0) == 1  
B.stride(1) == N
```

Location	Element
0	a_{11}
1	a_{21}
2	a_{12}
3	a_{22}

Row-Major AKA Right

- C++, NumPy (default)
- Rightmost extent is contiguous

Column-Major AKA Left

- Fortran, MATLAB
- Leftmost extent is contiguous

```
mdspan A{data, N, M};  
mdspan A{data, layout right::mapping{N, M}};
```

```
A[i, j]      == data[i * M + j]  
A.stride(0) == M  
A.stride(1) == 1
```

```
mdspan B{data, layout left::mapping{N, M}};
```

```
B[i, j]      == data[i + j * N]  
B.stride(0) == 1  
B.stride(1) == N
```

User-Defined Strides

```
mdspan C{data, layout stride::mapping{extents{N, M}, {X, Y}};
```

```
A[i, j]      == data[i * X + j * Y]  
A.stride(0) == X  
A.stride(1) == Y
```

Layouts map (i, j, k, \dots) to a storage location.

Layouts map (i, j, k, \dots) to a storage location.

Anyone can define a layout.

Layouts map (i, j, k, \dots) to a storage location.

Anyone can define a layout.

Layouts may:

➤ Be non-contiguous.

Layouts map (i, j, k, \dots) to a storage location.

Anyone can define a layout.

Layouts may:

- Be non-contiguous.
- Map multiple indices to the same location.

Layouts map (i, j, k, \dots) to a storage location.

Anyone can define a layout.

Layouts may:

- Be non-contiguous.
- Map multiple indices to the same location.
- Perform complicated computations.

Layouts map (i, j, k, \dots) to a storage location.

Anyone can define a layout.

Layouts may:

- Be non-contiguous.
- Map multiple indices to the same location.
- Perform complicated computations.
- Have or refer to state.

Parametric layouts enables
generic multi-dimensional algorithms.

```
void your_function(Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic>& m);
```

```
void your_function(Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic>& m);  
your_function(Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic>{...});
```

```
void your_function(Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic>& m);
```

```
your_function(Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic>{...});  
your_function(boost::numeric::ublas::matrix<double>{...});  
your_function(Mat{...}); // PETSc  
your_function(blaze::DynamicMatrix<double, blaze::rowMajor>{...});  
your_function(cutlass::HostTensor<float, cutlass::layout::ColumnMajor>{...});  
// ...
```

```
void your_function(std::mdspan<T, Extents, Layout, Accessor> m);  
  
your_function(Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic>{...});  
your_function(boost::numeric::ublas::matrix<double>{...});  
your_function(Mat{...}); // PETSc  
your_function(blaze::DynamicMatrix<double, blaze::rowMajor>{...});  
your_function(cutlass::HostTensor<float, cutlass::layout::ColumnMajor>{...});  
// ...
```

```

struct my_matrix {
public:
    my_matrix(std::size_t N, std::size_t M)
        : num_rows_(N), num_cols_(M), storage_(num_rows_ * num_cols_) {}

    double& operator()(size_t i, size_t j)
    { return storage_[i * num_cols_ + j]; }
    const double& operator()(size_t i, size_t j) const
    { return storage_[i * num_cols_ + j]; }

    std::size_t num_rows() const { return num_rows_; }
    std::size_t num_cols() const { return num_cols_; }

private:
    std::size_t          num_rows_, num_cols_;
    std::vector<double> storage_;
};

```



```

struct my_matrix {
public:
    my_matrix(std::size_t N, std::size_t M)
        : num_rows_(N), num_cols_(M), storage_(num_rows_ * num_cols_) {}

    double& operator()(size_t i, size_t j)
    { return storage_[i * num_cols_ + j]; }
    const double& operator()(size_t i, size_t j) const
    { return storage_[i * num_cols_ + j]; }

    std::size_t num_rows() const { return num_rows_; }
    std::size_t num_cols() const { return num_cols_; }

    operator std::mdspan<double, std::dextents<2>>() const
{ return {storage_, num_rows_, num_cols_}; }

private:
    std::size_t          num_rows_, num_cols_;
    std::vector<double> storage_;
};

```

```
std::span A{input,  N * M};  
std::span B{output, M * N};  
  
auto v = stdv::cartesian_product(  
    stdv::iota(0, N),  
    stdv::iota(0, M));  
  
std::for_each(ex::par_unseq,  
    begin(v), end(v),  
    [=] (auto idx) {  
        auto [i, j] = idx;  
        B[i + j * N] = A[i * M + j];  
    });
```

```
std::mdspan A{input,  N, M};  
std::mdspan B{output, M, N};  
  
auto v = stdv::cartesian_product(  
    stdv::iota(0, A.extent(0)),  
    stdv::iota(0, A.extent(1)));  
  
std::for_each(ex::par_unseq,  
    begin(v), end(v),  
    [=] (auto idx) {  
        auto [i, j] = idx;  
        B[j, i] = A[i, j];  
    });
```

```
std::mdspan A{input, N, M, 0};
```

```
std::mdspan B{output, N, M, 0};
```

```
auto v = stdv::cartesian_product(  
    stdv::iota(1, A.extent(0) - 1),  
    stdv::iota(1, A.extent(1) - 1),  
    stdv::iota(1, A.extent(2) - 1));
```

```
std::for_each(ex::par_unseq,  
    begin(v), end(v),  
    [=] (auto idx) {  
        auto [i, j, k] = idx;  
        B[i, j, k] = ( A[i, j, k-1] +  
                      A[i-1, j, k] +  
                      A[i, j-1, k] + A[i, j, k] + A[i, j+1, k]  
                      + A[i+1, j, k]  
                      + A[i, j, k+1] ) / 7.0  
    });
```

```

std::mdspan A{input,
               std::layout_left::mapping{N, M, O}};
std::mdspan B{output,
               std::layout_left::mapping{N, M, O}};

```

```

auto v = stdv::cartesian_product(
    stdv::iota(1, A.extent(0) - 1),
    stdv::iota(1, A.extent(1) - 1),
    stdv::iota(1, A.extent(2) - 1));

std::for_each(ex::par_unseq,
    begin(v), end(v),
    [=] (auto idx) {
        auto [i, j, k] = idx;
        B[i, j, k] = ( A[i, j, k-1] +
                       A[i-1, j, k] +
                       A[i, j-1, k] + A[i, j, k] + A[i, j+1, k]
                       + A[i+1, j, k]
                       + A[i, j, k+1] ) / 7.0
    });

```

```
submdspan(mdspan<...> m, SliceSpecifiers... ss)  
-> mdspan<...>
```

submdspan(mdspan<...> m, SliceSpecifiers... ss)
-> mdspan<...>

Slice Specifier	Argument	Reduces Rank?
Single Index	Integral	

submdspan(mdspan<...> m, SliceSpecifiers... ss)
-> mdspan<...>

Slice Specifier	Argument	Reduces Rank?
Single Index	Integral	✓
Range of Indices	std::pair<Integral, Integral> std::tuple<Integral, Integral>	✗

submdspan(mdspan<...> m, SliceSpecifiers... ss)
-> mdspan<...>

Slice Specifier	Argument	Reduces Rank?
Single Index	Integral	✓
Range of Indices	std::pair<Integral, Integral> std::tuple<Integral, Integral>	✗
All Indices	std::full_extent	✗

```
std::mdspan m0{data, 64, 128, 32};  
  
auto m1 = std::submdspan(m0, std::tuple{15, 23},  
                           std::tuple{31, 39},  
                           std::tuple{ 7, 15});
```

```
std::mdspan m0{data, 64, 128, 32};  
  
auto m1 = std::submdspan(m0, std::tuple{15, 23},  
                           std::tuple{31, 39},  
                           std::tuple{ 7, 15});  
  
m1.rank() == 3
```

```
std::mdspan m0{data, 64, 128, 32};

auto m1 = std::submdspan(m0, std::tuple{15, 23},
                           std::tuple{31, 39},
                           std::tuple{ 7, 15});

m1.rank()      == 3
m1.extent(0)   == 8
m1.extent(1)   == 8
m1.extent(2)   == 8
```

```
std::mdspan m0{data, 64, 128, 32};

auto m1 = std::submdspan(m0, std::tuple{15, 23},
                          std::tuple{31, 39},
                          std::tuple{ 7, 15});

m1.rank()      == 3
m1.extent(0)   == 8
m1.extent(1)   == 8
m1.extent(2)   == 8
m1[i, j, k]    == m0[i + 15, j + 31, k + 7]
```

```

std::mdspan m0{data, 64, 128, 32};

auto m1 = std::submdspan(m0, std::tuple{15, 23},
                        std::tuple{31, 39},
                        std::tuple{ 7, 15});

m1.rank()      == 3
m1.extent(0)   == 8
m1.extent(1)   == 8
m1.extent(2)   == 8
m1[i, j, k]    == m0[i + 15, j + 31, k + 7]

auto m2 = std::submdspan(m0, 15,
                        std::full_extent,
                        31);

```

```

std::mdspan m0{data, 64, 128, 32};

auto m1 = std::submdspan(m0, std::tuple{15, 23},
                          std::tuple{31, 39},
                          std::tuple{ 7, 15});

m1.rank()      == 3
m1.extent(0)   == 8
m1.extent(1)   == 8
m1.extent(2)   == 8
m1[i, j, k]    == m0[i + 15, j + 31, k + 7]

auto m2 = std::submdspan(m0, 15,
                          std::full_extent,
                          31);

m2.rank()      == 1

```

```

std::mdspan m0{data, 64, 128, 32};

auto m1 = std::submdspan(m0, std::tuple{15, 23},
                        std::tuple{31, 39},
                        std::tuple{ 7, 15});

m1.rank()      == 3
m1.extent(0)   == 8
m1.extent(1)   == 8
m1.extent(2)   == 8
m1[i, j, k]    == m0[i + 15, j + 31, k + 7]

auto m2 = std::submdspan(m0, 15,
                        std::full_extent,
                        31);

m2.rank()      == 1
m2.extent(0)   == 128

```



```

std::mdspan m0{data, 64, 128, 32};

auto m1 = std::submdspan(m0, std::tuple{15, 23},
                          std::tuple{31, 39},
                          std::tuple{ 7, 15});

m1.rank()      == 3
m1.extent(0)   == 8
m1.extent(1)   == 8
m1.extent(2)   == 8
m1[i, j, k]    == m0[i + 15, j + 31, k + 7]

auto m2 = std::submdspan(m0, 15,
                          std::full_extent,
                          31);

m2.rank()      == 1
m2.extent(0)   == 128
m2[j]          == m0[15, j, 31]

```

```
std::mdspan A{input,  N,  M};  
std::mdspan B{output, M,  N};  
std::size_t T = ...;
```

```
std::mdspan A{input,  N, M};  
std::mdspan B{output, M, N};  
std::size_t T = ...;  
  
auto outer = stdv::cartesian_product(stdv::iota(0, (N + T - 1) / T),  
                                     stdv::iota(0, (M + T - 1) / T));
```

```

std::mdspan A{input,  N, M};
std::mdspan B{output, M, N};
std::size_t T = ...;

auto outer = stdv::cartesian_product(stdv::iota(0, (N + T - 1) / T),
                                     stdv::iota(0, (M + T - 1) / T));

std::for_each(ex::par_unseq, begin(outer), end(outer),
  [=] (auto tile) {
    auto [x, y] = tile;
    ...
  });

```

```

std::mdspan A{input,  N, M};
std::mdspan B{output, M, N};
std::size_t T = ...;

auto outer = std::cartesian_product(std::iota(0, (N + T - 1) / T),
                                     std::iota(0, (M + T - 1) / T));

std::for_each(ex::par_unseq, begin(outer), end(outer),
  [=] (auto tile) {
    auto [x, y] = tile;
    std::tuple selectN{T * x, std::min(T * (x + 1), N)};
    std::tuple selectM{T * y, std::min(T * (y + 1), M)};

    ...
  });

```

```

std::mdspan A{input,  N, M};
std::mdspan B{output, M, N};
std::size_t T = ...;

auto outer = stdv::cartesian_product(stdv::iota(0, (N + T - 1) / T),
                                     stdv::iota(0, (M + T - 1) / T));

std::for_each(ex::par_unseq, begin(outer), end(outer),
  [=] (auto tile) {
    auto [x, y] = tile;
    std::tuple selectN{T * x, std::min(T * (x + 1), N)};
    std::tuple selectM{T * y, std::min(T * (y + 1), M)};

    auto TA = std::submdspan(A, selectN, selectM);
    auto TB = std::submdspan(B, selectM, selectN);

    ...
  });

```

```

std::mdspan A{input,  N, M};
std::mdspan B{output, M, N};
std::size_t T = ...;

auto outer = stdv::cartesian_product(stdv::iota(0, (N + T - 1) / T),
                                     stdv::iota(0, (M + T - 1) / T));

std::for_each(ex::par_unseq, begin(outer), end(outer),
  [=] (auto tile) {
    auto [x, y] = tile;
    std::tuple selectN{T * x, std::min(T * (x + 1), N)};
    std::tuple selectM{T * y, std::min(T * (y + 1), M)};

    auto TA = std::submdspan(A, selectN, selectM);
    auto TB = std::submdspan(B, selectM, selectN);

    auto inner = stdv::cartesian_product(stdv::iota(0, TA.extent(0)),
                                          stdv::iota(0, TA.extent(1)));

    ...
  });

```

```

std::mdspan A{input,  N, M};
std::mdspan B{output, M, N};
std::size_t T = ...;

auto outer = stdv::cartesian_product(stdv::iota(0, (N + T - 1) / T),
                                     stdv::iota(0, (M + T - 1) / T));

std::for_each(ex::par_unseq, begin(outer), end(outer),
  [=] (auto tile) {
    auto [x, y] = tile;
    std::tuple selectN{T * x, std::min(T * (x + 1), N)};
    std::tuple selectM{T * y, std::min(T * (y + 1), M)};

    auto TA = std::submdspan(A, selectN, selectM);
    auto TB = std::submdspan(B, selectM, selectN);

    auto inner = stdv::cartesian_product(stdv::iota(0, TA.extent(0)),
                                          stdv::iota(0, TA.extent(1)));

    for (auto [i, j] : inner)
      TB[j, i] = TA[i, j];
  });

```




Standard Algorithms

adjacent_difference	is_sorted[_until]	rotate[_copy]
adjacent_find	lexicographical_compare	search[_n]
all_of	max_element	set_difference
any_of	merge	set_intersection
copy[_if _n]	min_element	set_symmetric_difference
count[_if]	minmax_element	set_union
equal	mismatch	sort
fill[_n]	move	stable_partition
find[_end _first_of _if _if_not]	none_of	stable_sort
for_each	nth_element	swap_ranges
generate[_n]	partial_sort[_copy]	transform
includes	partition[_copy]	uninitialized_copy[_n]
inplace_merge	remove[_copy _copy_if _if]	uninitialized_fill[_n]
is_heap[_until]	replace[_copy _copy_if _if]	unique
is_partitioned	reverse[_copy]	unique_copy

```
std::mdspan A{..., N, M};  
std::mdspan x{..., M};  
std::mdspan y{..., N};  
  
// y = 3.0 A x + 2.0 y  
std::matrix_vector_product(  
    ex::par_unseq,  
    std::scaled(3.0, A), x,  
    std::scaled(2.0, y), y);
```

```
std::mdspan A{..., N, M};
```

```
std::mdspan x{..., M};
```

```
std::mdspan y{..., N};
```

```
//  $y = 3.0 A x + 2.0 y$ 
```

```
std::matrix_vector_product(
```

```
    ex::par_unseq,
```

```
    std::scaled(3.0, A), x,
```

```
    std::scaled(2.0, y), y);
```

```
std::mdspan A{..., N, M};  
std::mdspan x{..., M};  
std::mdspan y{..., N};  
  
// y = 3.0 A x + 2.0 y  
std::matrix_vector_product(  
    ex::par_unseq,  
    std::scaled(3.0, A), x,  
    std::scaled(2.0, y), y);
```

```
std::mdspan A{..., N, M};  
std::mdspan x{..., M};  
std::mdspan b{..., N};
```

```
// Solve  $A x = b$  where  $A = U^T U$ 
```

```
// Solve  $U^T c = b$ , using  $x$  to store  $c$ 
```

```
std::triangular matrix vector solve(ex::par_unseq,  
                                     std::transposed(A),  
                                     std::upper_triangle, std::explicit_diagonal,  
                                     b, x);
```

```
// Solve  $U x = c$ , overwriting  $x$  with result
```

```
std::triangular matrix vector solve(ex::par_unseq,  
                                     A,  
                                     std::upper_triangle, std::explicit_diagonal,  
                                     x);
```

```
std::mdspan A{..., N, M};  
std::mdspan x{..., M};  
std::mdspan b{..., N};
```

```
// Solve  $A x = b$  where  $A = U^T U$ 
```

```
// Solve  $U^T c = b$ , using  $x$  to store  $c$ 
```

```
std::triangular matrix vector solve(ex::par_unseq,  
                                     std::transposed(A),  
                                     std::upper_triangle, std::explicit_diagonal,  
                                     b, x);
```

```
// Solve  $U x = c$ , overwriting  $x$  with result
```

```
std::triangular matrix vector solve(ex::par_unseq,  
                                     A,  
                                     std::upper_triangle, std::explicit_diagonal,  
                                     x);
```

`mdspan` doesn't provide ranges and iterators
that enumerate its elements.

mdspan doesn't provide ranges and iterators
that enumerate its elements.

Why not?

mdspan doesn't provide ranges and iterators
that enumerate its elements.

Why not?

Performance.

Why do we want ranges and iterators?

Why do we want ranges and iterators?

- Parameterization.

Why do we want ranges and iterators?

- Parameterization.
- Composability.

Why do we want ranges and iterators?

- Parameterization.
- Composability.

They enable generic programming.

```
void f(std::range auto r) {  
    for (auto x : r)  
        ...  
}
```

```
std::vector<T> v{...};
```

```
f(v);
```

```
void f(std::range auto r) {  
    for (auto x : r)  
        ...  
}
```

```
void f(std::range auto r) {  
    for (auto x : r)  
        ...  
}
```

```
std::vector<T> v{...};  
  
f(v);  
  
f(v | std::reverse);
```



```
void f(std::range auto r) {  
    for (auto x : r)  
        ...  
}
```

```
std::vector<T> v{...};  
  
f(v);  
  
f(v | std::reverse);  
  
f(v | std::stride(2));
```

```
void f(std::range auto r) {  
    for (auto x : r)  
        ...  
}
```

```
std::vector<T> v{...};
```

```
f(v);
```

```
f(v | std::reverse);
```

```
f(v | std::stride(2));
```

```
f(v | std::reverse  
    | std::stride(2));
```

```
void f(std::range auto r) {  
    for (auto x : r)  
        ...  
}
```

```
std::vector<T> v{...};  
  
f(v);  
  
f(v | std::reverse);  
  
f(v | std::stride(2));  
  
f(v | std::reverse  
    | std::stride(2));  
  
f(std::iota(0, N));
```

We want to write generic multi-dimensional code that is:


- Storage agnostic.
- Rank agnostic.
- Layout agnostic.
- Iteration pattern agnostic.
- Composable.

```
std::vector<T> v{...};
```

```
std::vector<T> v{...};
```

Forward


v



```
std::vector<T> v{...};
```


Forward

v



Reverse


v | reverse



```
std::vector<T> v{...};
```


Forward

v




Reverse

v | reverse



Spiral


v | spiral(2)




```
std::mdspan A{..., N, M};
```


```
std::mdspan A{..., N, M};
```

**Row-Major
AKA Right**




```
std::mdspan A{..., N, M};
```

Row-Major
AKA Right




Column-Major
AKA Left




```
std::mdspan A{..., N, M};
```


**Row-Major
AKA Right**



**Column-Major
AKA Left**



**Morton Order
AKA Z-Order Curve**



A range is one way to iterate a sequence.

There may be other ways to iterate that sequence; those would be different ranges.

Indexwise

Elementwise

Indexwise

- Iterate through positions in MD space.
- Often uses multiple elements.
- Examples:
 - Matrix multiplication.
 - Sum rows/columns.
 - Stencils.

Elementwise

Indexwise

- Iterate through positions in MD space.
- Often uses multiple elements.
- Examples:
 - Matrix multiplication.
 - Sum rows/columns.
 - Stencils.

Elementwise

- Iterate through elements.
- Position in MD space doesn't matter.
- Examples:
 - Multiply/add by scalar.
 - Count non-zeros.
 - Commutative reduction.

Multi-Dimensional
Index Space
 (i, j, k, \dots)

Linear Storage
Location Space
 x

Multi-Dimensional
Index Space
 (i, j, k, \dots)

MD to storage is
typically fast, but
loses information.



Linear Storage
Location Space
 x

Multi-Dimensional
Index Space
 (i, j, k, \dots)

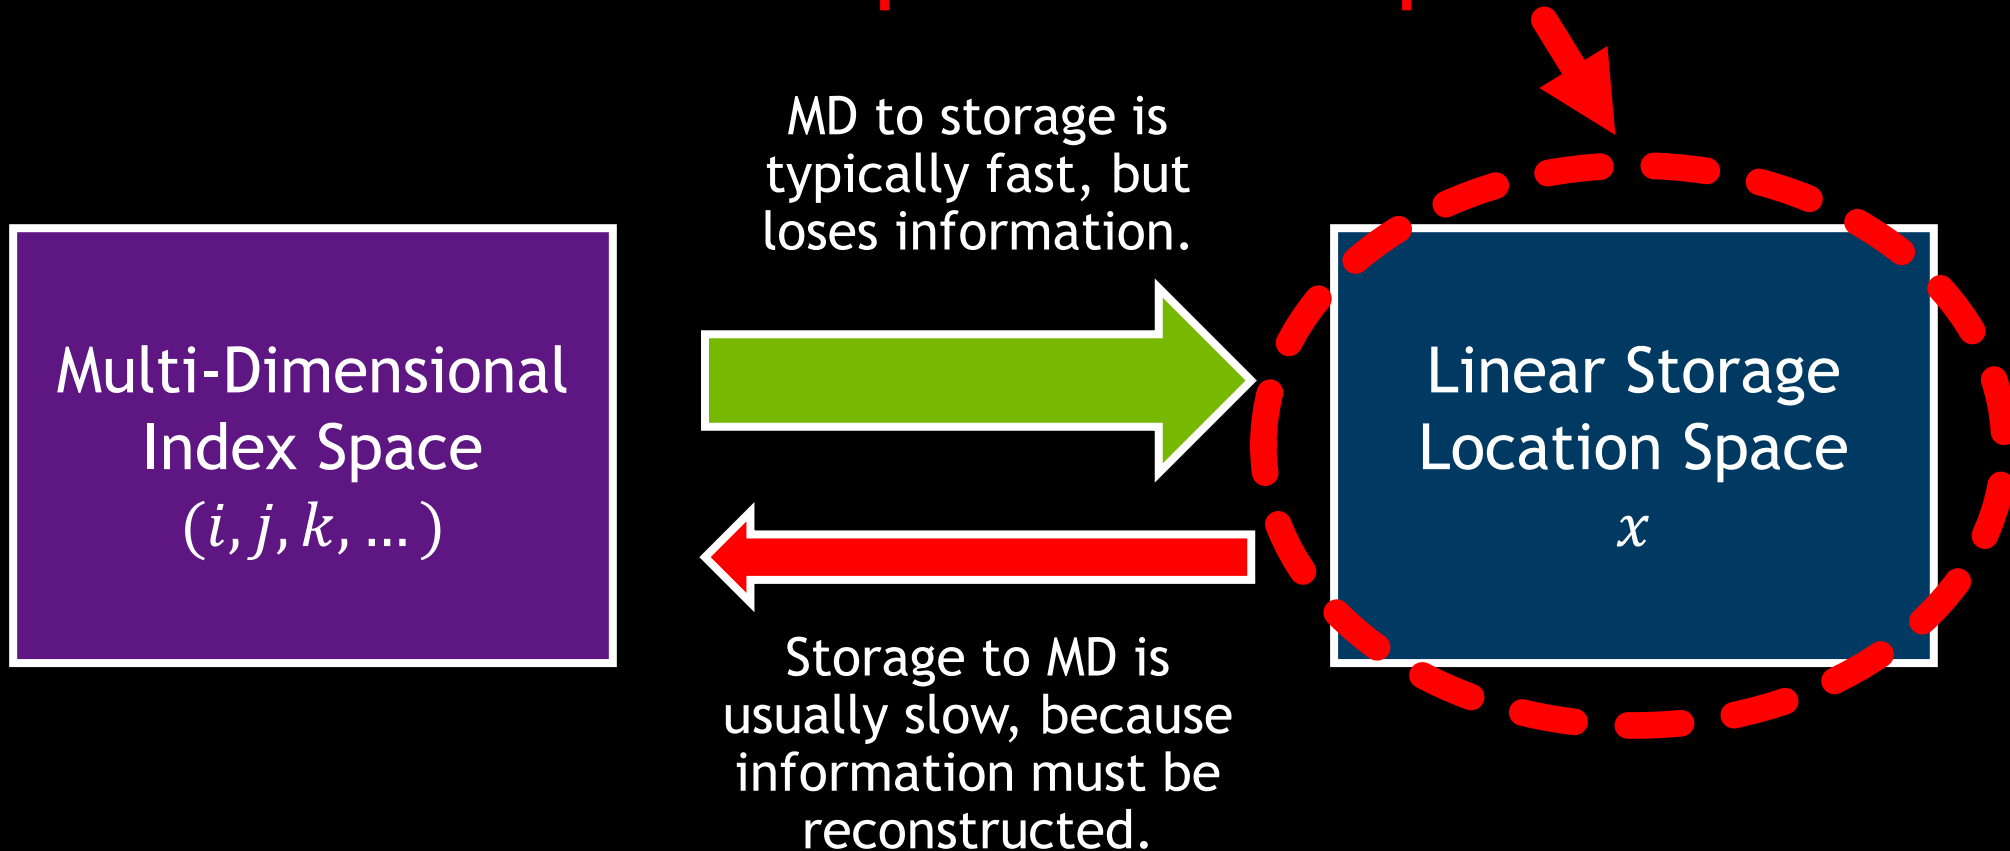
MD to storage is
typically fast, but
loses information.



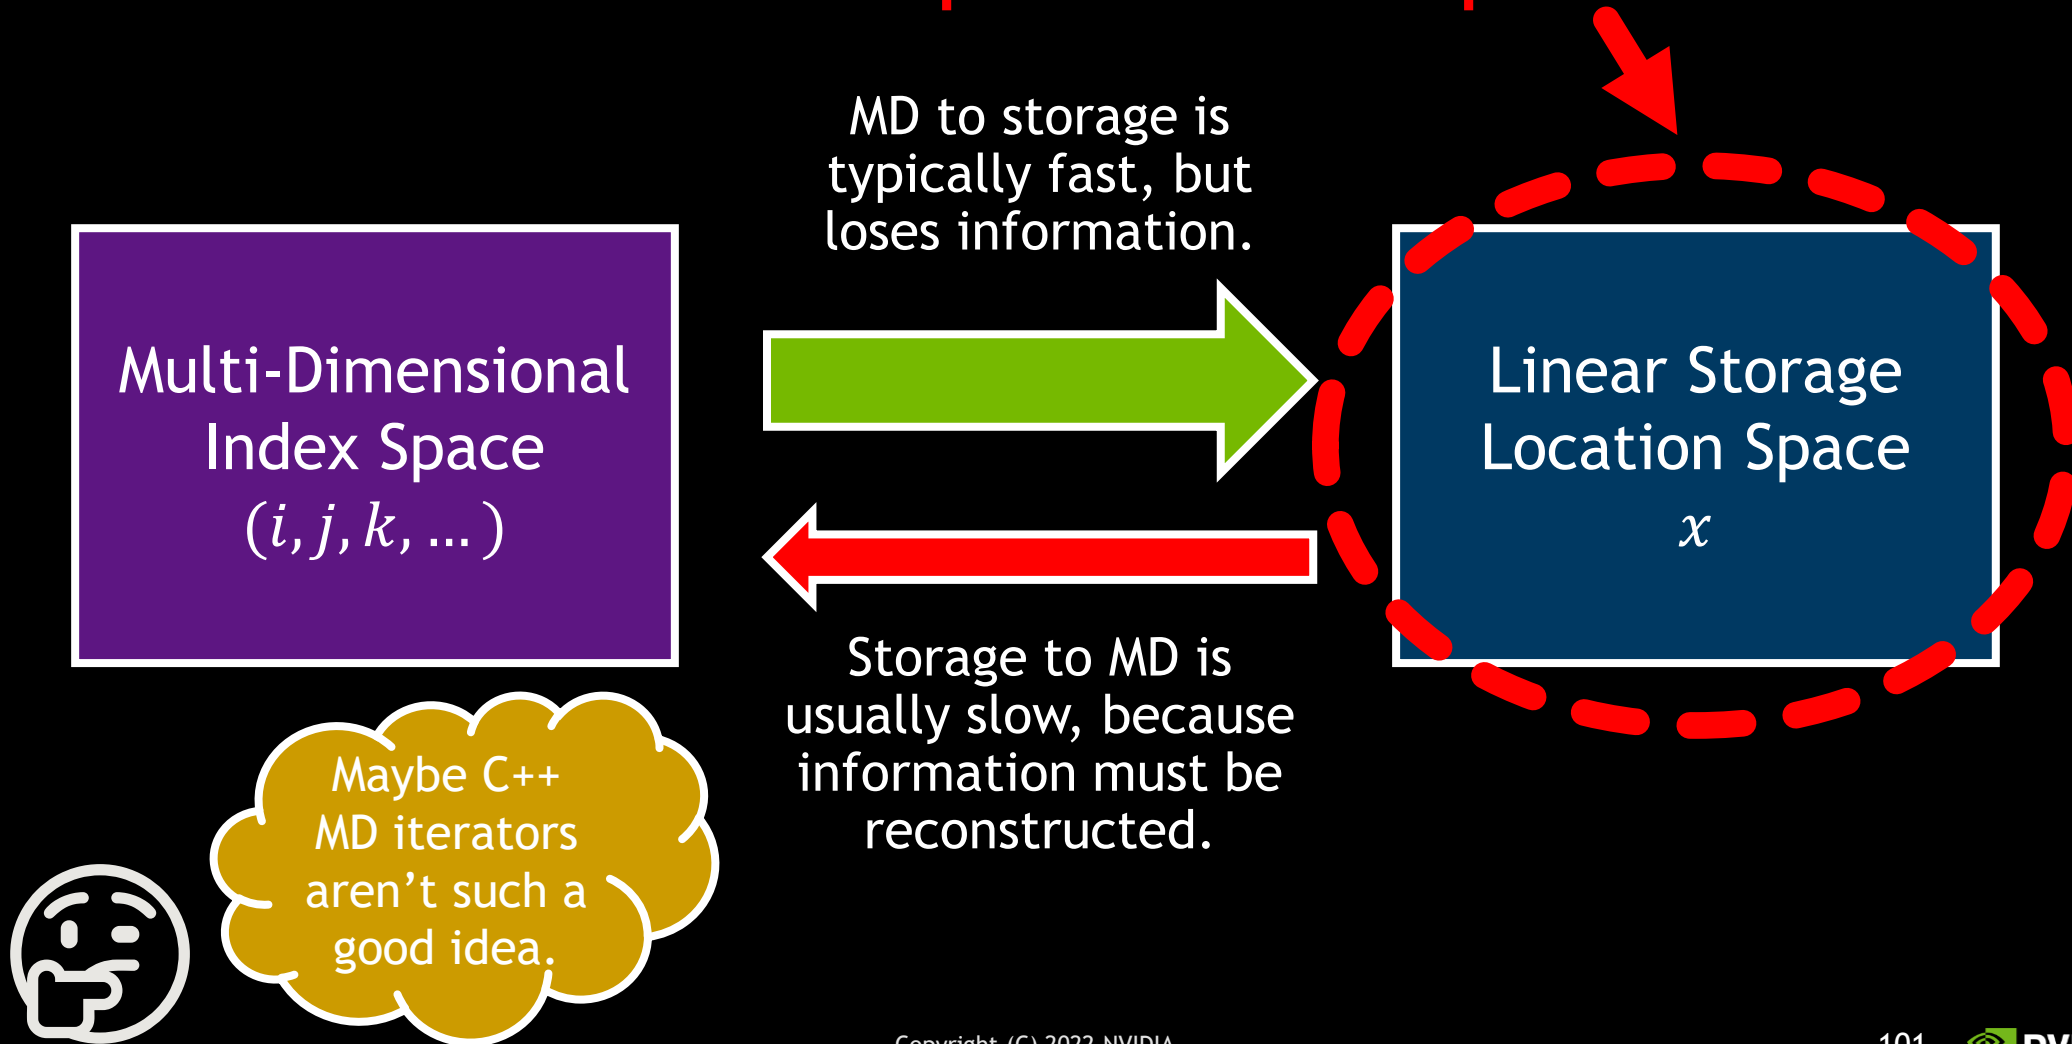
Linear Storage
Location Space
 x

Storage to MD is
usually slow, because
information must be
reconstructed.

C++ iterators operate in this space.



C++ iterators operate in this space.



```
for (int64_t j = 0; j != M; ++j)
    for (int64_t i = 0; i != N; ++i)
        A[i + j * N] = 0.0F;
```

```

for (int64_t j = 0; j != M; ++j)
    for (int64_t i = 0; i != N; ++i)
        A[i + j * N] = 0.0F;

```

```

pushq    %r15
pushq    %r14
pushq    %r12
pushq    %rbx
testq    %rsi, %rsi
je        .LBB0_12
testq    %rdi, %rdi
je        .LBB0_12
leaq     -32(%rdi), %r9
movq     %r9, %r8
shrq     $5, %r8
addq     $1, %r8
movl     %r8d, %r10d
andl     $3, %r10d
andq     $-4, %r8
leaq     960(%rdx), %rax
leaq     (,%rdi,8), %r14
leaq     192(%rdx), %r11
xorl     %r15d, %r15d
vxorps   %xmm0, %xmm0, %xmm0
xorl     %r12d, %r12d
jmp       .LBB0_3

.LBB0_11:
addq     $1, %r12
addq     %r14, %rax
addq     %rdi, %r15
addq     %r14, %rdx
cmpq     %rsi, %r12
je        .LBB0_12

.LBB0_3:
testq    %rdi, %rdi
movq     %rdx, %rcx
movq     %rdi, %rbx
je        .LBB0_13
cmpq     $96, %r9
jae       .LBB0_6
xorl     %ebx, %ebx
jmp       .LBB0_8

.LBB0_13:
movq     $0, (%rcx)
addq     $8, %rcx
addq     $-1, %rbx
jne       .LBB0_13
jmp       .LBB0_11

.LBB0_6:
movq     %r8, %rcx
xorl     %ebx, %ebx

.LBB0_7:
vmovups   %xmm0, -960(%rax,%rbx,8)
vmovups   %xmm0, -896(%rax,%rbx,8)
vmovups   %xmm0, -832(%rax,%rbx,8)
vmovups   %xmm0, -768(%rax,%rbx,8)
vmovups   %xmm0, -704(%rax,%rbx,8)
vmovups   %xmm0, -640(%rax,%rbx,8)
vmovups   %xmm0, -576(%rax,%rbx,8)
vmovups   %xmm0, -512(%rax,%rbx,8)
vmovups   %xmm0, -448(%rax,%rbx,8)
vmovups   %xmm0, -384(%rax,%rbx,8)
vmovups   %xmm0, -320(%rax,%rbx,8)
vmovups   %xmm0, -256(%rax,%rbx,8)
vmovups   %xmm0, -192(%rax,%rbx,8)
vmovups   %xmm0, -128(%rax,%rbx,8)
vmovups   %xmm0, -64(%rax,%rbx,8)
vmovups   %xmm0, (%rax,%rbx,8)
subq     $-128, %rbx
addq     $-4, %rcx
jne       .LBB0_7

.LBB0_8:
testq    %r10, %r10
je        .LBB0_11
addq     %r15, %rbx
leaq     (%r11,%rbx,8), %rcx
movq     %r10, %rbx

.LBB0_10:
vmovups   %xmm0, -192(%rcx)
vmovups   %xmm0, -128(%rcx)
vmovups   %xmm0, -64(%rcx)
vmovups   %xmm0, (%rcx)
addq     $256, %rcx
addq     $-1, %rbx
jne       .LBB0_10
jmp       .LBB0_11

.LBB0_12:
popq     %rbx
popq     %r12
popq     %r14
popq     %r15
retq

```

Setup

Outer Loop (Scalar)

Primary Inner Loop (Vectorized)
Width: 16 x 512 bit = 8192 bit

Remainder Inner Loop (Vectorized)
Width: 4 x 512 bit = 2048 bit

Cleanup

Peel Loop

Scalar.

Attempts to align and dispatch
to a better loop.

Primary Loop

Vectorized.

Assumes certain alignment and
of items remaining.

Remainder Loop(s)

Scalar or vectorized.

May assume certain alignment
and # of items remaining.


```
auto r = stdv::cartesian_product(  
    stdv::iota(0, N),  
    stdv::iota(0, M));  
for (auto [i, j] : r)  
    A[i, j] = 0.0F;
```

```

auto r = stdv::cartesian_product(
    stdv::iota(0, N),
    stdv::iota(0, M));
for (auto [i, j] : r)
    A[i, j] = 0.0F;

```

```

xorl    %r8d, %r8d
testq   %rsi, %rsi
movl    %edi, %ecx
cmovnel %r8d, %ecx
movslq  %ecx, %r10
cmpq    %rdi, %r10
je      .LBB4_3
movq    (%rdx), %r9
xorl    %edx, %edx
.LBB4_2:
movslq  %edx, %rax
movq    %rax, %rdx
imulq   %rdi, %rdx
addq    %r10, %rdx
movq    $0, (%r9,%rdx,8)
addl    $1, %eax
movslq  %eax, %rdx
xorl    %eax, %eax
cmpq    %rsi, %rdx
cmovel  %r8d, %edx
sete    %al
addl    %eax, %ecx
movslq  %ecx, %r10
cmpq    %rdi, %r10
jne     .LBB4_2
.LBB4_3:
retq

```

```

struct index_2d_iterator { // Column-Major AKA Left
private:
    std::array<std::int64_t, 2> indices;
    std::array<std::int64_t, 2> extents;

public:
    index_2d_iterator(index_2d_iterator const&);
    index_2d_iterator& operator=(index_2d_iterator const&);
    bool operator<=>(index_2d_iterator const&) const;

    index_2d_iterator& operator++() {
        ++indices[0]; // Inner loop iteration-expression.

        if (extents[0] == indices[0]) { // Inner loop condition.
            ++indices[1]; // Outer loop increment.
            indices[0] = 0; // Inner loop init-statement.
        }
        return *this;
    }

    std::array<std::int64_t, 2> operator*() const { return indices; }
};

```

```
for (auto [i, j] : index_2d_range(N, M))  
    A[i, j] = 0.0F;
```

```

for (auto [i, j] : index_2d_range(N, M))
    A[i, j] = 0.0F;

```

```

        movq    (%rdx), %r8
        xorl    %eax, %eax
        xorl    %edx, %edx
.LBB2_1:
        movq    %rax, %rcx
        imulq   %rdi, %rcx
        addq    %rdx, %rcx
        movq    $0, (%r8,%rcx,8)
        addq    $1, %rdx
        xorl    %r9d, %r9d
        cmpq    %rdi, %rdx
        sete    %r9b
        movl    $0, %ecx
        cmoveq   %rcx, %rdx
        addq    %r9, %rax
        cmpq    %rsi, %rax
        jne     .LBB2_1
        retq

```

```

struct storage_2d_iterator { // Column-Major AKA Left
private:
    std::int64_t          location;
    std::array<std::int64_t, 2> extents;

public:
    storage_2d_iterator(storage_2d_iterator const&);
    storage_2d_iterator& operator=(storage_2d_iterator const&);
    bool operator<=>(storage_2d_iterator const&) const;

    storage_2d_iterator& operator++() {
        ++location;
        return *this;
    }

    std::array<std::int64_t, 2> operator*() const {
        return {location % extents[1], location / extents[1]};
    }
};

```

```
for (auto [i, j] : storage_2d_range(N, M))  
    A[i, j] = 0.0F;
```

```

pushq   %rbp
pushq   %r15
pushq   %r14
pushq   %r13
pushq   %r12
pushq   %rbx
movq    %rsi, %rbx
imulq   %rdi, %rbx
testq   %rbx, %rbx
je      .LBB5_3
movq    (%rdx), %rax
movq    %rax, -24(%rsp)
vpbroadcastq %rsi, %xmm0
vpbroadcastq %rdi, %xmm1
addq    $-8, %rbx
shrq    $3, %rbx
addq    $1, %rbx
vmovdqud .LCPI5_0(%rip), %xmm2
vextracti32x4 $3, %xmm0, %xmm3
vpextrq $1, %xmm3, -32(%rsp)
vmovq    %xmm3, -40(%rsp)
vextracti32x4 $2, %xmm0, %xmm3
vpextrq $1, %xmm3, -48(%rsp)
vmovq    %xmm3, -56(%rsp)
vpsrlq   $32, %xmm0, %xmm12
vpsrlq   $32, %xmm1, %xmm4
vxorpd   %xmm9, %xmm0, %xmm0
vpbroadcastq .LCPI5_1(%rip), %xmm10
vpbroadcastq .LCPI5_2(%rip), %xmm11
movq    -48(%rsp), %r14
.LBB5_2:
vextracti32x4 $3, %xmm2, %xmm5
vpextrq $1, %xmm5, %rax
cqto
movq    -32(%rsp), %r8
idivq    %r9
movq    %rax, -16(%rsp)
vmovq    %xmm5, %rax
cqto
movq    -40(%rsp), %r9
idivq    %r9
movq    %rax, %rbp
vextracti32x4 $2, %xmm2, %xmm5
vpextrq $1, %xmm5, %rax
cqto
idivq    %r14
movq    %rax, %r10
vmovq    %xmm5, %rax
cqto
idivq    -56(%rsp)
movq    %rax, %r11
vextracti128 $1, %ymm2, %ymm5
vpextrq $1, %xmm5, %rax
vextracti128 $1, %ymm0, %xmm6
vpextrq $1, %xmm6, %rcx
movq    %rcx, -8(%rsp)
cqto
idivq    %rcx
movq    %rax, %rsi
vmovq    %xmm5, %rax
vmovq    %xmm6, %r15
cqto
idivq    %r15
movq    %rax, %rdi
vpextrq $1, %xmm2, %rax
vpextrq $1, %xmm0, %r12
cqto
idivq    %r12
movq    %rax, %rcx
vmovq    %xmm2, %rax
vmovq    %xmm0, %r13
cqto
idivq    %r13
vmovq    -16(%rsp), %xmm5
vmovq    %rbp, %xmm6
vpunpckldq %xmm5, %xmm6, %xmm5
vmovq    %r10, %xmm6
vmovq    %r11, %xmm7
vpunpckldq %xmm6, %xmm7, %xmm6
vinserti128 $1, %xmm5, %ymm6, %ymm5
vmovq    %rsi, %xmm6
vmovq    %rdi, %xmm7
vpunpckldq %xmm6, %xmm7, %xmm6
vmovq    %rcx, %xmm7
vmovq    %rax, %xmm3
vpunpckldq %xmm7, %xmm3, %xmm3
vinserti128 $1, %xmm6, %ymm3, %ymm3
vinserti64x4 $1, %ymm5, %zmm3, %zmm3
vpmuludq %xmm12, %zmm3, %zmm5
vpsrlq   $32, %zmm3, %zmm6
vpmuludq %zmm0, %zmm6, %zmm7
vpaddq   %zmm7, %zmm5, %zmm5

```

```

vpmuludq %zmm0, %zmm3, %zmm7
vpmuludq %zmm4, %zmm3, %zmm8
vpmuludq %zmm1, %zmm6, %zmm6
vpsllq   $32, %zmm5, %zmm5
vpaddq   %zmm6, %zmm8, %zmm6
vpsllq   $32, %zmm6, %zmm6
vpmuludq %zmm1, %zmm3, %zmm3
vpaddq   %zmm5, %zmm7, %zmm5
vpaddq   %zmm6, %zmm3, %zmm3
kxnorw   %k0, %k0, %k1
vpaddq   %zmm10, %zmm2, %zmm8
vpsubq   %zmm5, %zmm2, %zmm5
vextracti32x4 $3, %zmm8, %xmm6
vpextrq $1, %xmm6, %rax
vpaddq   %zmm5, %zmm3, %zmm3
movq    -24(%rsp), %rbp
vscatterqpd %zmm9, (%rbp,%zmm3,8) (%k1)
cqto
idivq    %r8
movq    %rax, %r8
vmovq    %xmm6, %rax
cqto
idivq    %r9
movq    %rax, %r9
vextracti32x4 $2, %zmm8, %xmm3
vpextrq $1, %xmm3, %rax
cqto
idivq    %r14
movq    %rax, %r10
vmovq    %xmm3, %rax
cqto
idivq    -56(%rsp)
movq    %rax, %r11
vextracti128 $1, %ymm8, %xmm3
vpextrq $1, %xmm3, %rax
cqto
idivq    -8(%rsp)
movq    %rax, %rsi
vmovq    %xmm3, %rax
cqto
idivq    %r15
movq    %rax, %rdi
vpextrq $1, %xmm8, %rax
cqto
idivq    %r12
movq    %rax, %rcx
vmovq    %xmm8, %rax
cqto
idivq    %r13
vmovq    %r8, %xmm3
vmovq    %r9, %xmm5
vpunpckldq %xmm3, %xmm5, %xmm3
vmovq    %r10, %xmm5
vmovq    %r11, %xmm6
vpunpckldq %xmm5, %xmm6, %xmm5
vinserti128 $1, %xmm3, %ymm5, %ymm5
vinserti64x4 $1, %ymm3, %zmm5, %zmm3
vpmuludq %zmm12, %zmm3, %zmm5
vpsrlq   $32, %zmm3, %zmm6
vpmuludq %zmm0, %zmm6, %zmm7
vpaddq   %zmm7, %zmm5, %zmm5
vpsllq   $32, %zmm5, %zmm5
vpmuludq %zmm5, %zmm7, %zmm7
vpaddq   %zmm5, %zmm3, %zmm7
vpmuludq %zmm1, %zmm6, %zmm6
vpaddq   %zmm6, %zmm7, %zmm6
vpsubq   %zmm5, %zmm7, %zmm5
vpsllq   $32, %zmm6, %zmm6
vpmuludq %zmm1, %zmm3, %zmm3
vpaddq   %zmm6, %zmm3, %zmm3
vpaddq   %zmm5, %zmm3, %zmm3
kxnorw   %k0, %k0, %k1
vpaddq   %zmm1, %zmm2, %zmm2
addq    $-2, %rbx
vscatterqpd %zmm9, (%rbp,%zmm3,8) (%k1)
jne      .LBB5_2
.LBB5_3:
popq     %rbx
popq     %r12
popq     %r13
popq     %r14
popq     %r15
popq     %rbp
retq

```

```

for (auto [i, j] : storage_2d_range(N, M))
    A[i, j] = 0.0F;

```



```

vextracti32x4 $3, %zmm2, %xmm5
vpextrq $1, %xmm5, %rax
cqto
movq -32(%rsp), %r8
idivq %r9
movq %rax, -16(%rsp)
vmovq %xmm5, %rax
cqto
movq -40(%rsp), %r9
idivq %r9
movq %rax, %rbp
vextracti32x4 $2, %zmm2, %xmm5
vpextrq $1, %xmm5, %rax
cqto
idivq %r14
movq %rax, %r10
vmovq %xmm5, %rax
cqto
idivq -56(%rsp)
movq %rax, %r11
vextracti128 $1, %ymm2, %xmm5
vpextrq $1, %xmm5, %rax
vextracti128 $1, %ymm0, %xmm6
vpextrq $1, %xmm6, %rcx
movq %rcx, -8(%rsp)
cqto
idivq %rcx
movq %rax, %rsi
vmovq %xmm5, %rax
vmovq %xmm6, %r15
cqto
idivq %r15
movq %rax, %rdi
vpextrq $1, %xmm2, %rax
vpextrq $1, %xmm0, %r12
cqto
idivq %r12
movq %rax, %rcx
vmovq %xmm2, %rax
vmovq %xmm0, %r13
cqto
idivq %r13
vmovq %rbp, %xmm6
vpunpcklqdq %xmm5, %xmm6, %xmm5
vmovq %r10, %xmm6
vmovq %r11, %xmm7
vpunpcklqdq %xmm6, %xmm7, %xmm6
vinseriti128 $1, %xmm5, %ymm6, %ymm5
vmovq %rsi, %xmm6
vmovq %rdi, %xmm7
vpunpcklqdq %xmm6, %xmm7, %xmm6
vmovq %rcx, %xmm7
vmovq %rax, %xmm3
vpunpcklqdq %xmm7, %xmm3, %xmm3
vinseriti128 $1, %xmm6, %ymm3, %ymm3
vinseriti64x4 $1, %ymm5, %zmm3, %zmm3
vpmuludq %zmm12, %zmm3, %zmm5
vpsrlq $32, %zmm3, %zmm6
vpmuludq %zmm0, %zmm6, %zmm7
vpaddq %zmm7, %zmm5, %zmm5
vpsllq $32, %zmm5, %zmm5
vpmuludq %zmm0, %zmm3, %zmm7
vpaddq %zmm5, %zmm7, %zmm5
vpmuludq %zmm4, %zmm3, %zmm7
vpmuludq %zmm1, %zmm6, %zmm6
vpaddq %zmm6, %zmm7, %zmm6
vpsubq %zmm5, %zmm8, %zmm5
vpsllq $32, %zmm6, %zmm6
vpmuludq %zmm1, %zmm3, %zmm3
vpaddq %zmm6, %zmm3, %zmm3
vpaddq %zmm5, %zmm3, %zmm3
kxnorw %k0, %k0, %k1
vpaddq %zmm11, %zmm2, %zmm2
addq $-2, %rbx
vscatterqpd %zmm9, (%rbp,%zmm3,8) {%k1}
jne .LBB5_2

vpextrq $1, %xmm6, %rax
vpaddq %zmm5, %zmm3, %zmm3
movq -24(%rsp), %rbp
vscatterqpd %zmm9, (%rbp,%zmm3,8) {%k1}
cqto
idivq %r8
movq %rax, %r8
vmovq %xmm6, %rax
cqto
idivq %r9
movq %rax, %r9
vextracti32x4 $2, %zmm8, %xmm3
vpextrq $1, %xmm3, %rax
cqto
idivq %r14
movq %rax, %r10
vmovq %xmm3, %rax
cqto
idivq -56(%rsp)
movq %rax, %r11
vextracti128 $1, %ymm8, %xmm3
vpextrq $1, %xmm3, %rax
cqto
idivq -8(%rsp)
movq %rax, %rsi
vmovq %xmm3, %rax
cqto
idivq %r15
movq %rax, %rdi
vpextrq $1, %xmm8, %rax
cqto
idivq %r12
movq %rax, %rcx
vmovq %xmm8, %rax
cqto
idivq %r13
vmovq %r8, %xmm3
vmovq %r9, %xmm5
vpunpcklqdq %xmm3, %xmm5, %xmm3
vmovq %r10, %xmm5
vmovq %r11, %xmm6
vpunpcklqdq %xmm5, %xmm6, %xmm5
vinseriti128 $1, %xmm3, %ymm5, %ymm3
vmovq %rsi, %xmm5
vmovq %rdi, %xmm6
vpunpcklqdq %xmm5, %xmm6, %xmm5
vmovq %rcx, %xmm6
vmovq %rax, %xmm7
vpunpcklqdq %xmm6, %xmm7, %xmm6
vinseriti128 $1, %xmm5, %ymm6, %ymm5
vinseriti64x4 $1, %ymm3, %zmm5, %zmm3
vpmuludq %zmm12, %zmm3, %zmm5
vpsrlq $32, %zmm3, %zmm6
vpmuludq %zmm0, %zmm6, %zmm7
vpaddq %zmm7, %zmm5, %zmm5
vpsllq $32, %zmm5, %zmm5
vpmuludq %zmm0, %zmm3, %zmm7
vpaddq %zmm5, %zmm7, %zmm5
vpmuludq %zmm4, %zmm3, %zmm7
vpmuludq %zmm1, %zmm6, %zmm6
vpaddq %zmm6, %zmm7, %zmm6
vpsubq %zmm5, %zmm8, %zmm5
vpsllq $32, %zmm6, %zmm6
vpmuludq %zmm1, %zmm3, %zmm3
vpaddq %zmm6, %zmm3, %zmm3
vpaddq %zmm5, %zmm3, %zmm3
kxnorw %k0, %k0, %k1
vpaddq %zmm11, %zmm2, %zmm2
addq $-2, %rbx
vscatterqpd %zmm9, (%rbp,%zmm3,8) {%k1}
jne .LBB5_2

```

```

vextracti32x4 $3, %zmm2, %xmm5
vpextrq $1, %xmm5, %rax
cqto
movq -32(%rsp), %r8
idivq %r9
movq %rax, -16(%rsp)
vmovq %xmm5, %rax
cqto
movq -40(%rsp), %r9
idivq %r9
movq %rax, %rbp
vextracti32x4 $2, %zmm2, %xmm5
vpextrq $1, %xmm5, %rax
cqto
idivq %r14
movq %rax, %r10
vmovq %xmm5, %rax
cqto
idivq -56(%rsp)
movq %rax, %r11
vextracti128 $1, %ymm2, %xmm5
vpextrq $1, %xmm5, %rax
vextracti128 $1, %ymm0, %xmm6
vpextrq $1, %xmm6, %rcx
movq %rcx, -8(%rsp)
cqto
idivq %rcx
movq %rax, %rsi
vmovq %xmm5, %rax
vmovq %xmm6, %r15
cqto
idivq %r15
movq %rax, %rdi
vpextrq $1, %xmm2, %rax
vpextrq $1, %xmm0, %r12
cqto
idivq %r12

movq %rax, %rcx
vmovq %xmm2, %rax
vmovq %xmm0, %r13
cqto
idivq %r13
vmovq -16(%rsp), %xmm5
vmovq %rbp, %xmm6
vpunpcklqdq %xmm5, %xmm6, %xmm5
vmovq %r10, %xmm6
vmovq %r11, %xmm7
vpunpcklqdq %xmm6, %xmm7, %xmm6
vinseriti128 $1, %xmm5, %ymm6, %ymm5
vmovq %rsi, %xmm6
vmovq %rdi, %xmm7
vpunpcklqdq %xmm6, %xmm7, %xmm6
vmovq %rcx, %xmm7
vmovq %rax, %xmm3
vpunpcklqdq %xmm7, %xmm3, %xmm3
vinseriti128 $1, %xmm6, %ymm3, %ymm3
vinseriti64x4 $1, %ymm5, %zmm3, %zmm3
vpmuludq %zmm12, %zmm3, %zmm5
vpsrlq $32, %zmm3, %zmm6
vpmuludq %zmm0, %zmm6, %zmm7
vpaddq %zmm7, %zmm5, %zmm5
vpsllq $32, %zmm5, %zmm5
vpmuludq %zmm0, %zmm3, %zmm7
vpaddq %zmm5, %zmm7, %zmm5
vpmuludq %zmm4, %zmm3, %zmm7
vpmuludq %zmm1, %zmm6, %zmm6
vpaddq %zmm6, %zmm7, %zmm6
vpsubq %zmm5, %zmm8, %zmm5
vpsllq $32, %zmm6, %zmm6
vpmuludq %zmm1, %zmm3, %zmm3
vpaddq %zmm6, %zmm3, %zmm3
vpaddq %zmm5, %zmm3, %zmm3
kxnorw %k0, %k0, %k1
vpaddq %zmm11, %zmm2, %zmm2
addq $-2, %rbx
vscatterqpd %zmm9, (%rbp,%zmm3,8) {%k1}
jne .LBB5_2

vpextrq $1, %xmm6, %rax
vpaddq %zmm5, %zmm3, %zmm3
movq -24(%rsp), %rbp
vscatterqpd %zmm9, (%rbp,%zmm3,8) {%k1}
cqto
idivq %r8
movq %rax, %r8
vmovq %xmm6, %rax
cqto
idivq %r9
movq %rax, %r9
vextracti32x4 $2, %zmm8, %xmm3
vpextrq $1, %xmm3, %rax
cqto
idivq %r14
movq %rax, %r10
vmovq %xmm3, %rax
cqto
idivq -56(%rsp)
movq %rax, %r11
vextracti128 $1, %ymm8, %xmm3
vpextrq $1, %xmm3, %rax
cqto
idivq -8(%rsp)
movq %rax, %rsi
vmovq %xmm3, %rax
cqto
idivq %r15
movq %rax, %rdi
vpextrq $1, %xmm8, %rax
cqto
idivq %r12
movq %rax, %rcx
vmovq %xmm8, %rax
cqto
idivq %r13
vmovq %r8, %xmm3

vmovq %r9, %xmm5
vpunpcklqdq %xmm3, %xmm5, %xmm3
vmovq %r10, %xmm5
vmovq %r11, %xmm6
vpunpcklqdq %xmm5, %xmm6, %xmm5
vinseriti128 $1, %xmm3, %ymm5, %ymm3
vmovq %rsi, %xmm5
vmovq %rdi, %xmm6
vpunpcklqdq %xmm5, %xmm6, %xmm5
vmovq %rcx, %xmm6
vmovq %rax, %xmm7
vpunpcklqdq %xmm6, %xmm7, %xmm6
vinseriti128 $1, %xmm5, %ymm6, %ymm5
vinseriti64x4 $1, %ymm3, %zmm5, %zmm3
vpmuludq %zmm12, %zmm3, %zmm5
vpsrlq $32, %zmm3, %zmm6
vpmuludq %zmm0, %zmm6, %zmm7
vpaddq %zmm7, %zmm5, %zmm5
vpsllq $32, %zmm5, %zmm5
vpmuludq %zmm0, %zmm3, %zmm7
vpaddq %zmm5, %zmm7, %zmm5
vpmuludq %zmm4, %zmm3, %zmm7
vpmuludq %zmm1, %zmm6, %zmm6
vpaddq %zmm6, %zmm7, %zmm6
vpsubq %zmm5, %zmm8, %zmm5
vpsllq $32, %zmm6, %zmm6
vpmuludq %zmm1, %zmm3, %zmm3
vpaddq %zmm6, %zmm3, %zmm3
vpaddq %zmm5, %zmm3, %zmm3
kxnorw %k0, %k0, %k1
vpaddq %zmm11, %zmm2, %zmm2
addq $-2, %rbx
vscatterqpd %zmm9, (%rbp,%zmm3,8) {%k1}
jne .LBB5_2

```

```
struct index_2d_generator;

index_2d_generator
generate_indices(std::int64_t N, std::int64_t M) {
    for (std::int64_t j = 0; j != M; ++j)
        for (std::int64_t i = 0; i != N; ++i)
            co_yield std::array{i, j};
}
```

```
for (auto [i, j] : generate_indices(N, M))  
    A[i, j] = 0.0F;
```

```

for (auto [i, j] : generate_indices(N, M))
    A[i, j] = 0.0F;

```

```

        movq    (%rdx), %r9
        addq    $56, %r9
        leaq    (,%rdi,8), %r8
        xorl    %edx, %edx
        vxorps  %xmm0, %xmm0, %xmm0
.LBB7_1:
        movq    %rdi, %rcx
        movq    %r9, %rax
.LBB7_2:
        vmovups %zmm0, -56(%rax)
        addq    $64, %rax
        addq    $-8, %rcx
        jne     .LBB7_2
        addq    $1, %rdx
        addq    %r8, %r9
        cmpq    %rsi, %rdx
        jne     .LBB7_1
        retq

```

Solutions

- Teach compilers to recognize non-loop control flow patterns and turn them into loops.

Solutions

- Teach compilers to recognize non-loop control flow patterns and turn them into loops.
- Develop compiler builtins that MD iterator authors can use to describe the inner loop.

Solutions

- Teach compilers to recognize non-loop control flow patterns and turn them into loops.
- Develop compiler builtins that MD iterator authors can use to describe the inner loop.
- Provide elementwise ranges and iterators that don't expose MD position.

Solutions

- Teach compilers to recognize non-loop control flow patterns and turn them into loops.
- Develop compiler builtins that MD iterator authors can use to describe the inner loop.
- Provide elementwise ranges and iterators that don't expose MD position.
- Don't use ranges and iterators for MD; build new constructs instead.

The range-based for statement:

for (for-range-declaration: for-range-initializer) statement

is equivalent to:

```
{  
    auto&& range = for-range-initializer;  
    auto begin = std::begin(range);  
    auto end = std::end(range);  
    for (; begin != end; ++begin) {  
        for-range-declaration = *begin;  
        statement  
    }  
}
```

[[stmt.ranged](#)]

The range-based for statement:

for (for-range-declaration: for-range-initializer) statement

[stmt.ranged]

The Range Protocol

```
template <typename Range>  
auto begin(Range&& range);
```

```
template <typename Range>  
auto end(Range&& range);
```

The Space Protocol

```
template <std::int64_t N, typename Space, typename... Outer>  
auto mbegin(Space&& space, std::tuple<Outer...>&& outer);
```

```
template <std::int64_t N, typename Space, typename... Outer>  
auto mdend(Space&& space, std::tuple<Outer...>&& outer);
```

```
template <typename Space>  
constexpr std::int64_t mdrank;
```

The Space Protocol

```
template <std::int64_t N, typename Space, typename... Outer>  
auto mbegin(Space&& space, std::tuple<Outer...>&& outer);
```

```
template <std::int64_t N, typename Space, typename... Outer>  
auto mend(Space&& space, std::tuple<Outer...>&& outer);
```

```
template <typename Space>  
constexpr std::int64_t mdrank;
```

The Space Protocol

```
template <std::int64_t N, typename Space, typename... Outer>  
auto mbegin(Space&& space, std::tuple<Outer...>&& outer);
```

```
template <std::int64_t N, typename Space, typename... Outer>  
auto mdend(Space&& space, std::tuple<Outer...>&& outer);
```

```
template <typename Space>  
constexpr std::int64_t mdrank;
```


The Space Protocol

```
template <std::int64_t N, typename Space, typename... Outer>  
auto mbegin(Space&& space, std::tuple<Outer...>&& outer);
```

```
template <std::int64_t N, typename Space, typename... Outer>  
auto mdend(Space&& space, std::tuple<Outer...>&& outer);
```

```
template <typename Space>  
constexpr std::int64_t mdrank;
```

The function:

```
template <typename Space, typename UnaryFunction>  
void mdfor(Space&& space, UnaryFunction&& f);
```

is equivalent to:

```
constexpr auto N = mdrank<MDSpace>;  
for (auto k = mdbegin<N - 1>(space); k != mdend<N - 1>(space); ++k)  
    for (auto j = mdbegin<N - 2>(space, *k); j != mdend<N - 2>(space, *k); ++j)  
        ...  
            for (auto i = mdbegin<0>(space, ...); i != mdend<0>(space, ...); ++i)  
                f(*i);
```

The space-based for statement:

for (for-space-declaration: for-space-initializer) statement

is equivalent to:

```
auto&& space = for-space-initializer;  
constexpr auto N = mdrank<MDSpace>;  
for (auto k = mdbegin<N - 1>(space); k != mdend<N - 1>(space); ++k)  
    for (auto j = mdbegin<N - 2>(space, *k); j != mdend<N - 2>(space, *k); ++j)  
        ...  
            for (auto i = mdbegin<0>(space, ...); i != mdend<0>(space, ...); ++i)  
                for-space-declaration = *i;  
                statement
```

```
std::mdspan A{..., N, M, O};
```

```
// Just the main diagonal.
```

```
A.indices() | stdv::filter([] (auto [i, j, k]) { return i == j && j == k; });
```

```
std::mdspan A{..., N, M, 0};
```

```
// Just the main diagonal.
```

```
A.indices() | stdv::filter([] (auto [i, j, k]) { return i == j && j == k; });
```

```
// Just [i, 0, 0].
```

```
A.indices() | on_extent<1>(stdv::filter([] (auto [j, k]) { return j == 0; }))  
| on_extent<2>(stdv::filter([] (auto [k]) { return k == 0; }));
```

```

std::mdspan A{..., N, M, 0};

// Just the main diagonal.
A.indices() | stdv::filter([] (auto [i, j, k]) { return i == j && j == k; });

// Just [i, 0, 0].
A.indices() | on_extent<1>(stdv::filter([] (auto [j, k]) { return j == 0; }))
              | on_extent<2>(stdv::filter([] (auto [k]) { return k == 0; }));

// Just interior points.
A.indices() | on_extent<0>(stdv::drop(1) | stdv::take(A.extent(0) - 2))
              | on_extent<1>(stdv::drop(1) | stdv::take(A.extent(1) - 2))
              | on_extent<2>(stdv::drop(1) | stdv::take(A.extent(2) - 2));

```

```
std::mdspan A{..., N, M};
```

```
// Traditional ranges & iterators for elementwise access.
```

```
// Multidimensional indices are NOT exposed from these.
```

```
std::random_access_range auto range = A;
```

```
std::random_access_iterator auto first = A.begin();
```

```
std::random_access_iterator auto last = A.end();
```

```
space auto indices = A.indices();
```

```
std::mdspan A{input,  N, M};  
std::mdspan B{output, M, N};  
  
auto v = stdv::cartesian_product(  
    stdv::iota(0, A.extent(0)),  
    stdv::iota(0, A.extent(1)));  
  
std::for_each(ex::par_unseq,  
    begin(v), end(v),  
    [=] (auto idx) {  
        auto [i, j] = idx;  
        B[j, i] = A[i, j];  
    });
```



```
std::mdspan A{input,  N,  M};  
std::mdspan B{output, M,  N};
```

```
stdr::for_each(  
    ex::par_unseq,  
    A.indices(),  
    [=] (auto [i, j]) {  
        B[j, i] = A[i, j];  
    });
```

```
std::mdspan A{input,  N, M};  
std::mdspan B{output, M, N};
```

```
ex::sender auto s =  
    ex::transfer_just(sch, A.indices())  
    | for_each_async(  
        [=] (auto [i, j]) {  
            B[j, i] = A[i, j];  
        });
```

```

std::mdspan A{input, N, M, 0};
std::mdspan B{output, N, M, 0};

auto v = stdv::cartesian_product(
    stdv::iota(1, A.extent(0) - 1),
    stdv::iota(1, A.extent(1) - 1),
    stdv::iota(1, A.extent(2) - 1));

std::for_each(ex::par_unseq,
    begin(v), end(v),
    [=] (auto idx) {
        auto [i, j, k] = idx;
        B[i, j, k] = ( A[i, j, k-1] +
                       A[i-1, j, k] +
                       A[i, j-1, k] + A[i, j, k] + A[i, j+1, k]
                       + A[i+1, j, k]
                       + A[i, j, k+1] ) / 7.0
    });

```

```

std::mdspan A{input,  N, M, O};
std::mdspan B{output, N, M, O};

std::for_each(ex::par_unseq,
  A.indices(),
  [=] (auto [i, j, k]) {
    B[i, j, k] = ( A[i, j, k-1] +
                  A[i-1, j, k] +
                  A[i, j-1, k] + A[i, j, k] + A[i, j+1, k]
                  + A[i+1, j, k]
                  + A[i, j, k+1] ) / 7.0
  });

```

```

std::mdspan A{input,  N, M, 0};
std::mdspan B{output, N, M, 0};

// Just interior points.
std::for_each(ex::par_unseq,
    A.indices() | on_extent<2>(stdv::drop(1) | stdv::take(A.extent(2)-2))
                | on_extent<1>(stdv::drop(1) | stdv::take(A.extent(1)-2))
                | on_extent<0>(stdv::drop(1) | stdv::take(A.extent(0)-2)),
    [=] (auto [i, j, k]) {
        B[i, j, k] = ( A[i, j, k-1] +
                      A[i-1, j, k] +
                      A[i, j-1, k] + A[i, j, k] + A[i, j+1, k]
                      + A[i+1, j, k]
                      + A[i, j, k+1] ) / 7.0
    });

```

```
std::mdspan A{input,  N, M, 0};  
std::mdspan B{output, N, M, 0};
```

```
// Just interior points.
```

```
for (auto [i, j, k] : A.indices()  
      | on_extent<0>(stdv::drop(1) | stdv::take(A.extent(0)-2))  
      | on_extent<1>(stdv::drop(1) | stdv::take(A.extent(1)-2))  
      | on_extent<2>(stdv::drop(1) | stdv::take(A.extent(2)-2)))  
  B[i, j, k] = ( A[i, j, k-1] +  
                 A[i-1, j, k] +  
                 A[i, j-1, k] + A[i, j, k] + A[i, j+1, k]  
                 + A[i+1, j, k]  
                 + A[i, j, k+1] ) / 7.0;
```

```

std::mdspan A{input,  N, M};
std::mdspan B{output, M, N};
std::size_t T = ...;

auto outer = stdv::cartesian_product(stdv::iota(0, (N + T - 1) / T),
                                     stdv::iota(0, (M + T - 1) / T));

std::for_each(ex::par_unseq, begin(outer), end(outer),
  [=] (auto tile) {
    auto [x, y] = tile;
    std::tuple selectN{T * x, std::min(T * (x + 1), N)};
    std::tuple selectM{T * y, std::min(T * (y + 1), M)};

    auto TA = std::submdspan(A, selectN, selectM);
    auto TB = std::submdspan(B, selectM, selectN);

    auto inner = stdv::cartesian_product(stdv::iota(0, TA.extent(0)),
                                          stdv::iota(0, TA.extent(1)));

    for (auto [i, j] : inner)
      TB[j, i] = TA[i, j];
  });

```

```

std::mdspan A{input, N, M};
std::mdspan B{output, M, N};
std::size_t T = ...;

auto outer = A.indices() | on_extent<1>(stdv::stride(T))
                        | on_extent<0>(stdv::stride(T));

std::for_each(ex::par_unseq, outer,
  [=] (auto [x, y]) {
    for (auto [i, j] : mdspan{{T * x, std::min(T * (x + 1), N)},
                              {T * y, std::min(T * (y + 1), M)}})
      B[j, i] = A[i, j];
  });

```



```

std::mdspan A{input, N, M};
std::mdspan B{output, M, N};
std::size_t T = ...;

for (auto [x, y] : A.indices() | on_extent<1>(stdv::stride(T))
                               | on_extent<0>(stdv::stride(T)))
    for (auto [i, j] : mdspan{{T * x, std::min(T * (x + 1), N)},
                              {T * y, std::min(T * (y + 1), M)}})
        B[j, i] = A[i, j];

```



@blelbach

